

# feature-engineering

November 17, 2025

## 0.0.1 Feature Engineering

```
[1]: %load_ext autoreload  
%autoreload 2
```

```
[2]: from mldl_hw3.preprocessing import DataLoader  
from mldl_hw3.consts import Brand  
from mldl_hw3.experiment import Experiment, ExperimentConfig  
  
from typing import Optional  
  
from xgboost import XGBRegressor  
from sklearn.pipeline import Pipeline  
from sklearn.base import BaseEstimator, TransformerMixin  
from sklearn.compose import TransformedTargetRegressor  
from sklearn.preprocessing import OneHotEncoder, StandardScaler  
from sklearn.linear_model import LinearRegression  
from sklearn.cluster import KMeans  
from category_encoders import TargetEncoder  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns  
import numpy as np
```

```
[3]: df_train, df_test = DataLoader("../dataset").load()
```

```
[4]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Index: 4428 entries, G4XLU0 to J2RCU8  
Data columns (total 15 columns):  
 #   Column           Non-Null Count  Dtype     
 ---  --     
 0   Name            4428 non-null    category  
 1   Location        4428 non-null    category  
 2   Year            4428 non-null    Int64  
 3   Kilometers_Driven 4428 non-null    Int64  
 4   Fuel_Type        4428 non-null    category  
 5   Transmission     4428 non-null    category
```

```

6   Owner_Type          4428 non-null  category
7   Mileage              4428 non-null  float64
8   Engine                4428 non-null  Int64
9   Power                 4336 non-null  float64
10  Colour               4428 non-null  category
11  Seats                  4428 non-null  Int64
12  Doors                  4428 non-null  Int64
13  New_Price             601 non-null  float64
14  Price                 4428 non-null  float64
dtypes: Int64(5), category(6), float64(4)
memory usage: 408.6+ KB

```

```
[5]: X_train = df_train.copy()
y_train = X_train.pop("Price")
X_test = df_test.drop(columns=["Price"])
```

## 0.0.2 Baseline

```
[6]: class CategoricalEncoder(TransformerMixin, BaseEstimator):
    def fit(
        self, X: pd.DataFrame, y: Optional[pd.Series] = None
    ) -> "CategoricalEncoder":
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()

        for feature in X.select_dtypes(["category"]):
            X[feature] = X[feature].cat.codes

        return X
```

```
[7]: def build_baseline_pipeline(**model_params) -> Pipeline:
    return Pipeline(
        [
            ("encoder", CategoricalEncoder()),
            ("model", XGBRegressor(**model_params)),
        ]
    )
```

```
[8]: baseline_exp = Experiment(
    ExperimentConfig(name="baseline", pipeline=build_baseline_pipeline())
)

baseline_exp_result = baseline_exp.run(X_train, y_train, X_test)
```

```
[Experiment: baseline]
Cross-validating (5-folds)...
```

```

CV score: 0.1606 ± 0.0163
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/baseline.csv
Experiment complete

```

**Brand extraction from Name** From the analysis on `preprocessing.ipynb`, we found that `Name` feature has too many categories. The car's name usually consists of brand and model. From common knowledge, we know that the car's brand has significant impact on forming its price. Based on this domain knowledge and to reduce category, we extract `Brand` from `Name`.

```

[9]: class BrandModelExtractor(TransformerMixin, BaseEstimator):
    def __init__(self,
                 name_col: str = "Name",
                 brand_col: str = "Brand",
                 model_col: str = "Model",
                 brand_enum=Brand,
                 ):
        self.name_col = name_col
        self.brand_col = brand_col
        self.model_col = model_col
        self.brand_enum = brand_enum

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> Union["BrandModelExtractor", None]:
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()

        names = X[self.name_col].str.title()
        brands = pd.Series([None] * len(names), index=names.index, dtype=object)
        models = pd.Series([None] * len(names), index=names.index, dtype=object)

        for brand in Brand:
            condition = names.str.startswith(brand.value, na=False) & brands.isna()

            if condition.any():
                brands.loc[condition] = brand.name
                matched_names = names.loc[condition]
                residuals = matched_names.str[len(brand.value):].str.strip()
                models.loc[condition] = residuals.where(residuals != "", None)

        brands = brands.rename("Brand").astype("category")

```

```

models = models.rename("Model").astype("category")

X[self.brand_col] = brands
X[self.model_col] = models

return X

```

[10]:

```

X_train = df_train.copy()
y_train = X_train.pop("Price")

df_train_bm = pd.DataFrame(BrandModelExtractor().fit_transform(X_train, y_train))
df_train_bm

```

[10]:

ID	Name	Location	Year	Kilometers_Driven	Fuel_Type
G4XLU0	Tata Indigo	Coimbatore	2013	59138	Diesel
CRSHOS	Toyota Corolla	Kochi	2013	81504	Diesel
FUJ4X1	Ford Ikon	Hyderabad	2007	92000	Petrol
QMVK6E	Hyundai i20	Kolkata	2012	33249	Diesel
4SWHFC	Honda City	Bangalore	2011	65000	Petrol
...	...	...	...	...	...
TR7SLB	Mahindra XUV500	Kochi	2016	51884	Diesel
QB41QE	Honda Jazz	Kolkata	2016	27210	Diesel
ODG8N7	Land Rover Range	Pune	2015	52000	Diesel
EV2ZBX	Maruti Alto	Delhi	2013	56000	Petrol
J2RCU8	Mercedes-Benz GL-Class	Bangalore	2014	52000	Diesel

ID	Transmission	Owner_Type	Mileage	Engine	Power	Colour	Seats
G4XLU0	Manual	First	17.00	1405	70.00	Others	5
CRSHOS	Manual	First	21.43	1364	87.20	Others	5
FUJ4X1	Manual	First	13.80	1299	70.00	Others	5
QMVK6E	Manual	First	21.27	1396	88.76	Black/Silver	5
4SWHFC	Manual	First	17.00	1497	118.00	White	5
...	...	...	...	...	...	...	...
TR7SLB	Manual	First	16.00	2179	140.00	White	7
QB41QE	Manual	First	27.30	1498	98.60	Others	5
ODG8N7	Automatic	First	12.70	2179	187.70	White	5
EV2ZBX	Manual	First	24.70	796	47.30	Others	5
J2RCU8	Automatic	First	12.00	2987	224.00	Black/Silver	7

ID	Doors	New_Price	Brand	Model
G4XLU0	4	NaN	Tata	Indigo
CRSHOS	4	NaN	Toyota	Corolla
FUJ4X1	4	NaN	Ford	Ikon

```

QMVK6E      4      NaN      Hyundai      I20
4SWHFC      4      NaN      Honda       City
...
TR7SLB      5      ...      Mahindra    Xuv500
QB41QE      4      NaN      Honda       Jazz
ODG8N7      4      NaN      Land_Rover Range
EV2ZBX      4      NaN      Maruti     Alto
J2RCU8      5      NaN      Mercedes_Benz G1-Class

```

[4428 rows x 16 columns]

Note that Brand and Model has an hierarchical relationship.

```
[11]: display(
    df_train_bm[["Brand", "Model"]].groupby(["Brand", "Model"], observed=True).
    ↪count()
)
```

Empty DataFrame  
Columns: []  
Index: [(Audi, A3), (Audi, A4), (Audi, A6), (Audi, A7), (Audi, A8), (Audi, Q3),  
(Audi, Q5), (Audi, Q7), (Audi, Rs5), (Audi, Tt), (BMW, 1), (BMW, 3), (BMW, 5),  
(BMW, 6), (BMW, 7), (BMW, X1), (BMW, X3), (BMW, X5), (BMW, X6), (BMW, Z4),  
(Bentley, Continental), (Chevrolet, Aveo), (Chevrolet, Beat), (Chevrolet,  
Captiva), (Chevrolet, Cruze), (Chevrolet, Enjoy), (Chevrolet, Optra),  
(Chevrolet, Sail), (Chevrolet, Spark), (Datsun, Go), (Datsun, Redi), (Datsun,  
Redi-Go), (Fiat, Avventura), (Fiat, Grande), (Fiat, Linea), (Fiat, Petra),  
(Fiat, Punto), (Fiat, Siena), (Force, One), (Ford, Aspire), (Ford, Ecosport),  
(Ford, Endeavour), (Ford, Fiesta), (Ford, Figo), (Ford, Freestyle), (Ford,  
Fusion), (Ford, Ikon), (Honda, Accord), (Honda, Amaze), (Honda, Br-V), (Honda,  
Brio), (Honda, Brv), (Honda, City), (Honda, Civic), (Honda, Cr-V), (Honda,  
Jazz), (Honda, Mobilio), (Honda, Wr-V), (Honda, Wrv), (Hyundai, Accent),  
(Hyundai, Creta), (Hyundai, Elantra), (Hyundai, Elite), (Hyundai, Eon),  
(Hyundai, Getz), (Hyundai, Grand), (Hyundai, I10), (Hyundai, I20), (Hyundai,  
Santa), (Hyundai, Santro), (Hyundai, Sonata), (Hyundai, Tucson), (Hyundai,  
Verna), (Hyundai, Xcent), (Isuzu, D-Max), (Isuzu, Mux), (Jaguar, F), (Jaguar,  
Xe), (Jaguar, Xf), (Jaguar, Xj), (Jeep, Compass), (Lamborghini, Gallardo),  
(Land\_Rover, Discovery), (Land\_Rover, Freelander), (Land\_Rover, Range),  
(Mahindra, Bolero), (Mahindra, Jeep), (Mahindra, Kuv), (Mahindra, Logan),  
(Mahindra, Nuvosport), (Mahindra, Quanto), (Mahindra, Renault), (Mahindra,  
Scorpio), (Mahindra, Ssangyong), (Mahindra, Thar), (Mahindra, Tuv), (Mahindra,  
Verito), (Mahindra, Xuv300), (Mahindra, Xuv500), (Mahindra, Xylo), ...]

[204 rows x 0 columns]

```
[12]: df_train_bm.Brand.unique()
```

```
[12]: ['Tata', 'Toyota', 'Ford', 'Hyundai', 'Honda', ..., 'Smart', 'Isuzu', 'Volvo',
'Bentley', 'Lamborghini']
Length: 29
Categories (29, object): ['Audi', 'BMW', 'Bentley', 'Chevrolet', ..., 'Tata',
'Toyota', 'Volkswagen', 'Volvo']
```

```
[13]: df_train_bm.Brand.value_counts()
```

```
[13]: Brand
Maruti           884
Hyundai          817
Honda            462
Toyota           309
Mercedes_Benz    239
Volkswagen       225
Ford              217
Mahindra          199
BMW               191
Audi              177
Tata              140
Skoda             117
Renault            117
Chevrolet          82
Nissan             64
Land_Rover         43
Jaguar             28
Fiat               23
Mini               21
Mitsubishi         20
Porsche            14
Volvo              14
Jeep                11
Datsun              7
Isuzu              3
Lamborghini         1
Force               1
Smart               1
Bentley              1
Name: count, dtype: int64
```

```
[14]: df_train_bm.Model.unique()
```

```
[14]: ['Indigo', 'Corolla', 'Ikon', 'I20', 'City', ..., 'Cls-Class', 'Gallardo',
'Wr-V', 'Outlander', 'Xc90']
Length: 201
Categories (201, object): ['1', '3', '5', '6', ..., 'Yeti', 'Z4', 'Zen', 'Zest']
```

```
[15]: with pd.option_context("display.max_rows", None):
    display(df_train_bm.Model.value_counts())
```

Model	
Swift	253
City	206
I20	183
Verna	129
Grand	125
Innova	121
I10	120
Wagon	113
Alto	107
Polo	107
Xuv500	89
New	80
Amaze	80
Fortuner	77
Vento	76
3	75
Ecosport	72
Figo	71
Creta	69
Duster	65
E-Class	63
Ertiga	58
A4	55
Santro	53
Ciaz	51
Corolla	51
Etios	50
Ritz	48
5	47
Baleno	45
Brio	44
Jazz	44
Eon	44
Celerio	44
Scorpio	40
Xcent	39
A6	36
Vitara	35
Kwid	34
Superb	34
Indigo	29
Beat	29
Civic	27
Q7	27

Endeavour	27
Indica	27
Rapid	26
Q5	25
Fiesta	25
X1	24
Zen	24
Octavia	23
Micra	23
Range	23
Q3	22
Dzire	21
Accord	21
Laura	20
Sx4	20
Xf	20
Cr-V	19
Cooper	19
Sunny	19
Zest	19
M-Class	18
Nano	18
Terrano	18
Gla	17
Jetta	17
Ameo	17
X5	16
S	15
Xylo	15
Bolero	15
Ikon	15
Eeco	15
Pajero	15
Aveo	13
Elantra	13
Santa	13
A-Star	13
Cruze	13
Manza	12
Linea	12
Accent	12
Compass	11
Discovery	10
Mobilio	10
Freelander	10
Gl-Class	10
Gle	9
Fabia	9

Cla	9
X3	9
Spark	8
B	8
Optra	8
Tiago	8
Omni	8
Kuv	8
800	7
Cayenne	7
Ssangyong	7
Elite	7
Tuv	6
Getz	6
Camry	6
Glc	6
Safari	6
Sail	6
Pulse	5
Xj	5
6	5
Yeti	5
X6	5
Esteem	5
7	5
Passat	5
Grande	5
Scala	4
Sonata	4
S60	4
Brv	4
S-Class	4
Quanto	4
Panamera	4
Aspire	4
Ignis	4
Xc60	4
A	4
V40	4
Avventura	3
Thar	3
Tigor	3
Tucson	3
Bolt	3
Verito	3
A7	3
A3	3
Wrv	3

S-Cross	3
1	3
Redi-Go	3
Enjoy	3
Br-V	3
Fluence	3
R-Class	3
Captur	3
Qualis	3
Koleos	3
Go	3
Hexa	3
Renault	2
X-Trail	2
Logan	2
Lancer	2
A8	2
Jeep	2
Gls	2
Crosspolo	2
D-Max	2
Cayman	2
Xe	2
Xenon	2
Freestyle	2
Z4	2
Rs5	2
Tt	2
Captiva	2
Nuvosport	2
C-Class	2
Sumo	2
Fusion	1
Xc90	1
Beetle	1
Punto	1
Gallardo	1
Xuv300	1
Slc	1
F	1
Fortwo	1
S80	1
Boxster	1
Redi	1
Siena	1
Nexon	1
Slk-Class	1
Mux	1

```

Countryman      1
Platinum       1
Continental    1
Wr-V           1
Petra          1
Outlander      1
Clubman         1
One             1
Cls-Class      1
Cedia           1
Teana           1
Montero         1
Evalia          1
Name: count, dtype: int64

```

Model feature has too high cardinality, and there are non-negligible amount of categories with very small sample. Since we already have Brand, we could just drop Model. Alternatively, we could try target encoding.

0. Baseline (only Name)
1. Categorical Brand and Model added

```
[16]: extract_brand_model_exp = Experiment(
    ExperimentConfig(
        name="extract-brand-model",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

extract_brand_model_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```

[Experiment: extract-brand-model]
Cross-validating (5-folds)...
CV score: 0.1905 ± 0.0190
    +0.0299  +0.0027 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/extract-brand-model.csv
Experiment complete

```

2. Drop Model

```
[17]: class FeatureDropper(TransformerMixin, BaseEstimator):
    def __init__(self, cols: list[str]):
        self.cols = cols

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> "FeatureDropper":
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        return X.drop(columns=self.cols)
```

```
[18]: drop_model_exp = Experiment(
    ExperimentConfig(
        name="drop-model",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                ("drop_model", FeatureDropper(["Model"])),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

drop_model_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

[Experiment: drop-model]  
Cross-validating (5-folds)...  
CV score: 0.1606 ± 0.0163  
+0.0000 +0.0000 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/drop-model.csv  
Experiment complete

### 3. Drop Name

```
[19]: drop_name_exp = Experiment(
    ExperimentConfig(
        name="drop-name",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                ("drop_name", FeatureDropper(["Name"])),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        )
)
```

```

        ),
    )
)

drop_name_exp.run(X_train, y_train, X_test, baseline_exp_result);

```

[Experiment: drop-name]  
Cross-validating (5-folds)...  
CV score: 0.2379 ± 0.0157  
+0.0772 -0.0006 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/drop-name.csv  
Experiment complete

#### 4. Target encoding Model. (Leave Brand categorical)

```
[20]: model_te_exp = Experiment(
    ExperimentConfig(
        name="model-te",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                ("target_encode", TargetEncoder(cols=["Model"])),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

model_te_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

[Experiment: model-te]  
Cross-validating (5-folds)...  
CV score: 0.1523 ± 0.0107  
-0.0083 -0.0056 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/model-te.csv  
Experiment complete

#### 5. Target encoding Brand and Model.

```
[21]: brand_model_te_exp = Experiment(
    ExperimentConfig(
        name="brand-model-te",
```

```

        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                ("target_encode", TargetEncoder(cols=["Brand", "Model"])),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

brand_model_te_exp.run(X_train, y_train, X_test, baseline_exp_result);

```

[Experiment: brand-model-te]  
Cross-validating (5-folds)...  
CV score: 0.1496 ± 0.0170  
-0.0110 +0.0007 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/brand-model-te.csv  
Experiment complete

## 6. Target encoding Brand, Model, and Name.

```
[22]: brand_model_name_te_exp = Experiment(
    ExperimentConfig(
        name="brand-model-name-te",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name"]),
                ),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

brand_model_name_te_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

[Experiment: brand-model-name-te]  
Cross-validating (5-folds)...  
CV score: 0.1505 ± 0.0136  
-0.0101 -0.0027 compared to baseline (Negative is better)  
Training on full training set...

```
Creating submission on test set...
Submission created: artifacts/experiment-results/brand-model-name-te.csv
Experiment complete
```

Target encoding Brand, Model, and Name yields the biggest gain.

Thinking deeper on the hierarchical relationship of Brand and Model, the difference of Model target encoded value and Brand target encoded value can tell how high class is the car in that brand. Test whether this helps.

```
[23]: class ModelPremiumEncoder(TransformerMixin, BaseEstimator):
    def __init__(self,
                 brand_col: str = "Brand",
                 model_col: str = "Model",
                 premium_col: str = "Model_Premium",
                 ):
        self.brand_col = brand_col
        self.model_col = model_col
        self.premium_col = premium_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series] = None) -> "ModelPremiumEncoder":
        return self

    def transform(self, X: pd.DataFrame):
        X = X.copy()
        X[self.premium_col] = X[self.model_col] - X[self.brand_col]
        return X
```

```
[24]: X_train = df_train.copy()
y_train = X_train.pop("Price")
X_test = df_test.copy().drop(columns=["Price"])

brand_model_name_mp_te_exp = Experiment(
    ExperimentConfig(
        name="model-premium",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name"]),
                ),
                ("model_premium", ModelPremiumEncoder()),
                ("category_encode", CategoricalEncoder()),
            ]
        )
    )
)
```

```

        ("model", XGBRegressor()),
    ],
),
)

brand_model_name_mp_te_exp.run(X_train, y_train, X_test, baseline_exp_result);

```

```

[Experiment: model-premium]
Cross-validating (5-folds)...
CV score: 0.1485 ± 0.0161
    -0.0121 -0.0002 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/model-premium.csv
Experiment complete

```

Model\_Premium improved the performance slightly. However, the gain is marginal and standard deviation is worse than without it. I think this is not worth it and will stick with target encoded Brand, Model, and Name.

### Location

```
[25]: df_train.Location.value_counts()
```

```
[25]: Location
Mumbai      586
Hyderabad   546
Coimbatore  478
Kochi        475
Pune         465
Delhi        420
Kolkata     384
Chennai      362
Jaipur       315
Bangalore    237
Ahmedabad   160
Name: count, dtype: int64
```

There is no location with significantly small datapoints.

Since location can influence car price due to regional demand, wealth, urban/rural, taxes, ... target encoding could provide a valuable signal of this regional effect.

1. Categorical Location (Same as baseline)

```
[26]: baseline_exp.run(X_train, y_train);
```

```
[Experiment: baseline]
Cross-validating (5-folds)...
```

```
CV score: 0.1606 ± 0.0163
Training on full training set...
Experiment complete
```

### 1. Target encoded Location

```
[27]: X_train = df_train.copy()
y_train = X_train.pop("Price")
X_test = df_test.copy().drop(columns=["Price"])

location_te_exp = Experiment(
    ExperimentConfig(
        name="location target encoding",
        pipeline=Pipeline(
            [
                (
                    "target_encode",
                    TargetEncoder(cols=["Location"]),
                ),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

location_te_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```
[Experiment: location target encoding]
Cross-validating (5-folds)...
CV score: 0.1559 ± 0.0162
    -0.0047 -0.0001 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/location target encoding.csv
Experiment complete
```

It does help to target encode Location.

Combining with the Name feature engineering.

```
[28]: name_location_best_exp = Experiment(
    ExperimentConfig(
        name="name, location best feature engineering",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
```

```

        (
            "target_encode",
            TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
        ),
        ("category_encode", CategoricalEncoder()),
        ("model", XGBRegressor()),
    ]
),
)
)

name_location_best_exp.run(X_train, y_train, X_test, baseline_exp_result);

[Experiment: name, location best feature engineering]
Cross-validating (5-folds)...
CV score: 0.1480 ± 0.0155
-0.0127 -0.0008 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/name, location best feature
engineering.csv
Experiment complete

```

Combined with engineered Name, it further increased the performance.

- brand-model-name-te:  $0.1505 \pm 0.0136$
- brand-model-name-location-te:  $0.1480 \pm 0.0155$

### Year

```
[29]: df_train.Year.value_counts()
```

```
[29]: Year
2014    571
2016    564
2015    558
2013    484
2017    443
2012    411
2011    342
2010    247
2018    216
2009    146
2008    123
2007     93
2019     75
2006     55
2005     44
```

```

2004    21
2002    12
2003    11
2001     5
2000     3
1999     2
1998     2
Name: count, dtype: Int64

```

As explored in `preprocessing.ipynb`, calculating the `Age` might be more relevant in terms of price. Assuming that current year is 2020, I should try transforming `Year` to `Age`.

```
[30]: class YearToAgeTransformer(TransformerMixin, BaseEstimator):
    def __init__(self, year_col: str = "Year", age_col: str = "Age", current_year: int = 2020):
        self.year_col = year_col
        self.age_col = age_col
        self.current_year = current_year

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series] = None) -> "YearToAgeTransformer":
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()
        X[self.age_col] = self.current_year - X[self.year_col]
        X.drop(columns=[self.year_col], inplace=True)
        return X
```

0. Year (Same as baseline)

1. Year transformed to Age

```
[31]: year_to_age_transform_exp = Experiment(
    ExperimentConfig(
        name="year-to-age-transform",
        pipeline=Pipeline(
            [
                ("transform", YearToAgeTransformer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)
```

```

year_to_age_transform_exp.run(X_train, y_train, X_test, baseline_exp_result);

[Experiment: year-to-age-transform]
Cross-validating (5-folds)...
CV score: 0.1602 ± 0.0154
    -0.0004 -0.0009 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/year-to-age-transform.csv
Experiment complete

```

This transformation does not necessarily add more information but makes it more interpretable and nice to work with other features.

```

[32]: name_location_year_exp = Experiment(
    ExperimentConfig(
        name="name-location-year-engineered",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

name_location_year_exp.run(X_train, y_train, X_test, baseline_exp_result);

[Experiment: name-location-year-engineered]
Cross-validating (5-folds)...
CV score: 0.1479 ± 0.0159
    -0.0127 -0.0004 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/name-location-year-
engineered.csv
Experiment complete

```

Together with the Name and Location engineered, transforming Year to Age did not help at all. It might be due to Name -> Brand, Model, Name being the dominant feature. However, I will keep this

transformation since it makes the feature more interpretable while not hurting the performance. It could make this feature nicer to work with when analyzing interactions.

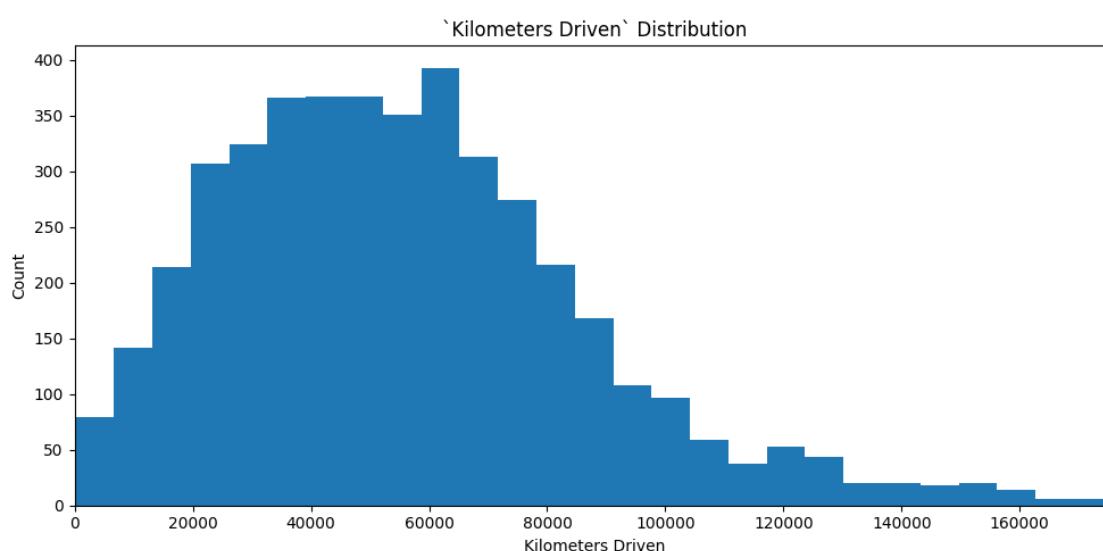
### Kilometers Driven

```
[33]: df_train.Kilometers_Driven
```

```
[33]: ID
G4XLU0      59138
CRSHOS      81504
FUJ4X1       92000
QMVK6E       33249
4SWHFC      65000
...
TR7SLB      51884
QB41QE      27210
ODG8N7       52000
EV2ZBX      56000
J2RCU8       52000
Name: Kilometers_Driven, Length: 4428, dtype: Int64
```

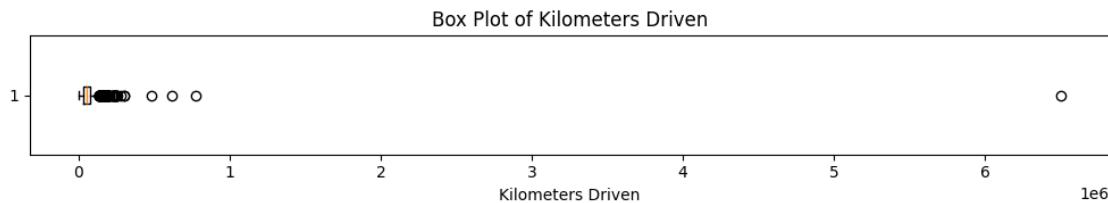
```
[34]: plt.figure(figsize=(10, 5))
plt.hist(df_train.Kilometers_Driven, bins=1000)
limit = np.percentile(df_train.Kilometers_Driven, 99)
plt.xlim(0, limit)
plt.xlabel("Kilometers Driven")
plt.ylabel("Count")
plt.title("`Kilometers Driven` Distribution")

plt.tight_layout()
plt.show()
```



This looks like a Gamma distribution. It makes sense because if we model the `Kilometers_Driven` as an accumulation of daily driving (with random noise added to consistent daily driving), it will yield a Gamma distribution.

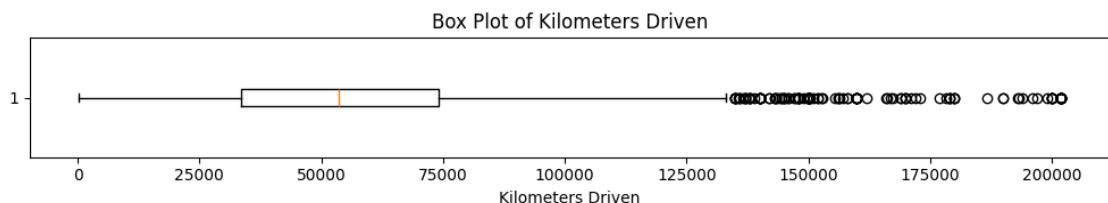
```
[35]: plt.figure(figsize=(10, 2))
plt.boxplot(df_train["Kilometers_Driven"], vert=False)
plt.xlabel("Kilometers Driven")
plt.title("Box Plot of Kilometers Driven")
plt.tight_layout()
plt.show()
```



There are clearly some extreme impossible data points. They should be clipped.

```
[36]: df_train.Kilometers_Driven = df_train.Kilometers_Driven.clip(
    upper=int(df_train.Kilometers_Driven.quantile(0.995))
)
df_test.Kilometers_Driven = df_test.Kilometers_Driven.clip(
    upper=int(df_test.Kilometers_Driven.quantile(0.995))
)
```

```
[37]: plt.figure(figsize=(10, 2))
plt.boxplot(df_train["Kilometers_Driven"], vert=False)
plt.xlabel("Kilometers Driven")
plt.title("Box Plot of Kilometers Driven")
plt.tight_layout()
plt.show()
```



```
[38]: class KilometersDrivenClipper(TransformerMixin, BaseEstimator):
    def __init__(
        self,
        kilometers_driven_col: str = "Kilometers_Driven",
        clipping_quantile: float = 0.995,
    ):
        self.kilometers_driven_col = kilometers_driven_col
        self.clipping_quantile = clipping_quantile

    def fit(
        self, X: pd.DataFrame, y: Optional[pd.Series] = None
    ) -> "KilometersDrivenClipper":
        return self

    def transform(self, X: pd.DataFrame):
        X = X.copy()
        X[self.kilometers_driven_col] = X[self.kilometers_driven_col].clip(
            upper=int(X[self.kilometers_driven_col].quantile(self.
            ↪clipping_quantile)))
        )
        return X
```

1. Kilometers\_Driven
1. Kilometers\_Driven Outliers clipped

```
[39]: kilometers_driven_outlier_clip_exp = Experiment(
    ExperimentConfig(
        name="clipping-extreme-Kilometers_Driven",
        pipeline=Pipeline(
            [
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

kilometers_driven_outlier_clip_exp.run(X_train, y_train, X_test, ↴
    baseline_exp_result);
```

[Experiment: clipping-extreme-Kilometers\_Driven]  
Cross-validating (5-folds)...  
CV score: 0.1596 ± 0.0165  
-0.0010 +0.0002 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/clipping-extreme-

```
Kilometers_Driven.csv
Experiment complete
```

The performance increased slightly but also increased std slightly. The effect is negligible but still good to have for numerical stability.

```
[40]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

current_best_exp.run(X_train, y_train, None, baseline_exp_result);

[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1485 ± 0.0167
-0.0121 +0.0004 compared to baseline (Negative is better)
Training on full training set...
Experiment complete
```

### Fuel\_Type

```
[41]: df_train.Fuel_Type.value_counts()
```

```
[41]: Fuel_Type
Diesel      2372
Petrol     2008
CNG         42
LPG          6
Electric      0
Name: count, dtype: int64
```

```
[42]: df_test.Fuel_Type.value_counts()
```

```
[42]: Fuel_Type
Diesel      790
Petrol      685
CNG         13
LPG          3
Electric     0
Name: count, dtype: int64
```

Fuel\_Type has low cardinality but CNG, LPG, and Electric has extremely low datapoints. Grouping CNG, LPG, and Electric to Other could mitigate this problem.

```
[43]: class FuelTypeGrouper(TransformerMixin, BaseEstimator):
    def __init__(self, fuel_type_col: str = "Fuel_Type"):
        self.fuel_type_col = fuel_type_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> "FuelTypeGrouper":
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()
        X[self.fuel_type_col] = (
            X[self.fuel_type_col]
            .astype("object")
            .replace({"CNG": "Other", "LPG": "Other", "Electric": "Other"})
            .astype("category")
        )
        return X
```

```
[44]: fuel_type_grouping_exp = Experiment(
    ExperimentConfig(
        name="grouping-infrequent-fuel-type",
        pipeline=Pipeline(
            [
                ("group_infrequent_fuel_type", FuelTypeGrouper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

fuel_type_grouping_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```
[Experiment: grouping-infrequent-fuel-type]
Cross-validating (5-folds)...
CV score: 0.1601 ± 0.0150
-0.0005 -0.0013 compared to baseline (Negative is better)
Training on full training set...
```

```
Creating submission on test set...
Submission created: artifacts/experiment-results/grouping-infrequent-fuel-
type.csv
Experiment complete
```

Due to low cardinality, (Petrol / Diesel / Other), I do not feel the need to engineer it further (e.g. target encoding).

```
[45]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("group_infrequent_fuel_type", FuelTypeGrouper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

current_best_exp.run(X_train, y_train, None, baseline_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1489 ± 0.0151
-0.0117 -0.0012 compared to baseline (Negative is better)
Training on full training set...
Experiment complete
```

## Transmission

```
[46]: df_train.Transmission.value_counts()
```

```
[46]: Transmission
      Manual      3162
      Automatic   1266
      Name: count, dtype: int64
```

Transmission has two categories and well balanced. There is no need to engineer this feature alone further.

## Owner Type

```
[47]: df_train.Owner_Type.cat.categories
```

```
[47]: Index(['First', 'Second', 'Third', 'Fourth & Above'], dtype='object')
```

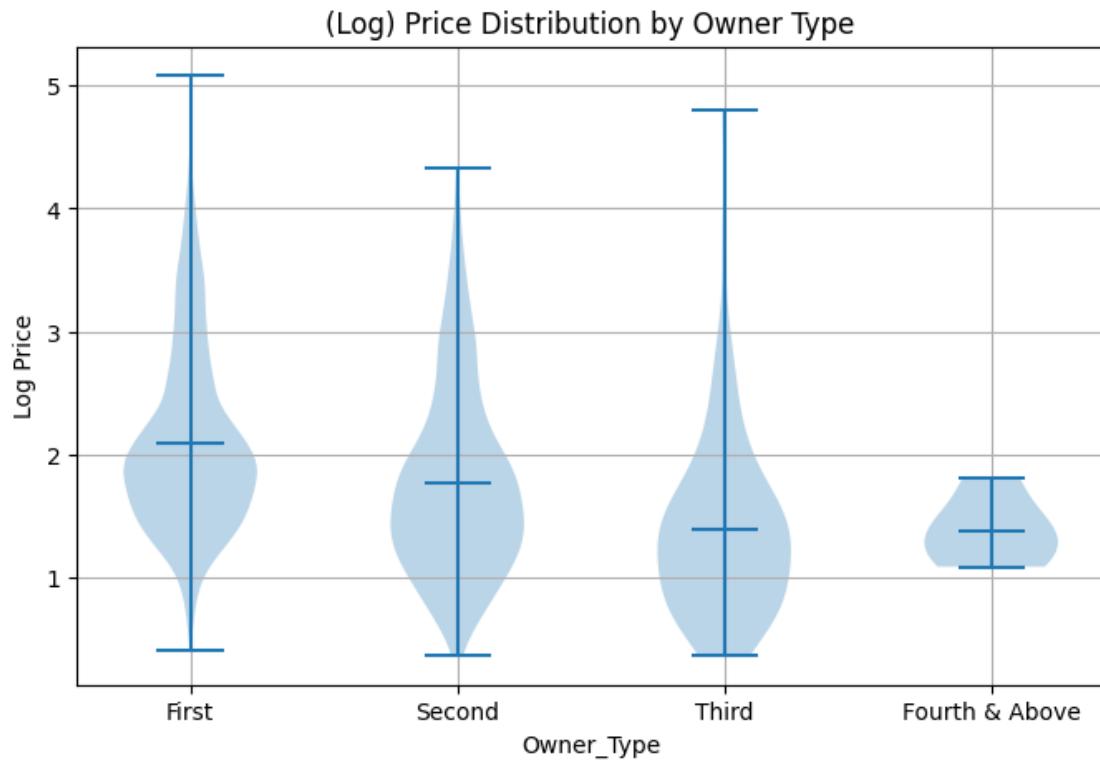
```
[48]: df_train.Owner_Type.value_counts()
```

```
[48]: Owner_Type
```

```
First           3617
Second          715
Third           91
Fourth & Above      5
Name: count, dtype: int64
```

Based on common knowledge, Owner\_Type might be one of the most significant predictor of Price.

```
[49]: categories = list(df_train.Owner_Type.cat.categories)
plt.figure(figsize=(8, 5))
plt.violinplot(
    [
        np.log1p(df_train[df_train.Owner_Type == category]["Price"])
        for category in categories
    ],
    showmeans=True,
)
plt.xticks(ticks=range(1, len(categories) + 1), labels=categories)
plt.xlabel("Owner_Type")
plt.ylabel("Log Price")
plt.title("(Log) Price Distribution by Owner Type")
plt.grid(True)
plt.show()
```

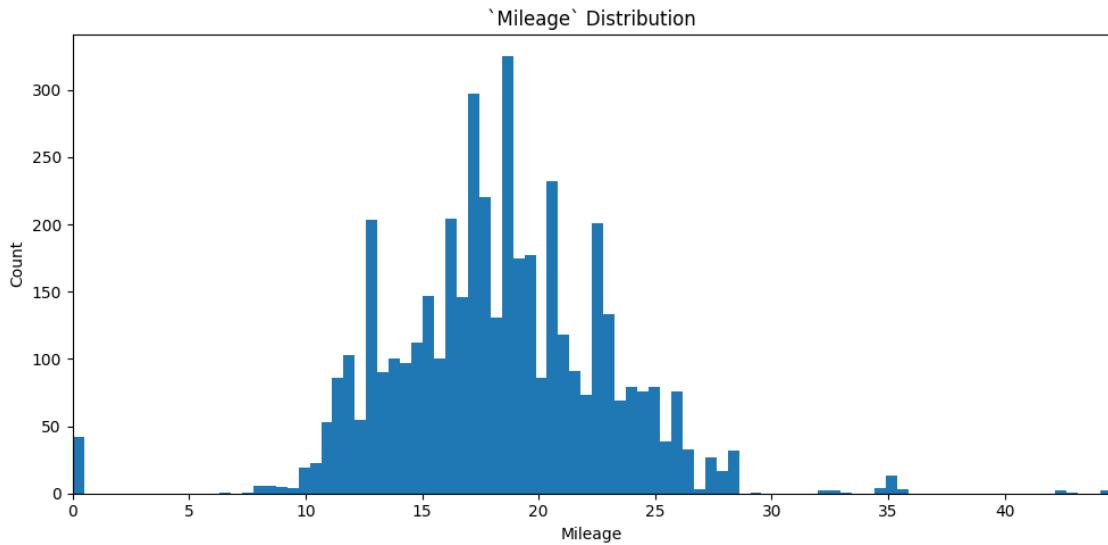


Clearly, as `Owner_Type` increases, `Price` drops monotonically. (`Owner_Type` is already encoded as ordinal category in preprocessing). It is good enough for now.

### Mileage

```
[50]: plt.figure(figsize=(10, 5))
plt.hist(df_train.Mileage, bins=100)
limit = np.percentile(df_train.Mileage, 99.9)
plt.xlim(0, limit)
plt.xlabel("Mileage")
plt.ylabel("Count")
plt.title("`Mileage` Distribution")

plt.tight_layout()
plt.show()
```



```
[51]: with pd.option_context("display.max_rows", None):
    display(df_train.Mileage.value_counts().sort_index())
```

Mileage	Count
0.0000	42
6.4000	1
7.5000	1
7.8100	1
7.9400	2
8.0000	1
8.2000	2
8.4500	1
8.5000	1
8.6000	2
8.7000	2
9.0000	5
9.4300	1
9.5000	2
9.5200	1
9.7400	2
9.8000	3
9.9000	2
10.0000	5
10.1000	4
10.1300	3
10.2000	3
10.3700	2
10.4000	2
10.5000	16

10.8000	7
10.9000	6
10.9100	10
10.9300	5
10.9800	1
11.0000	14
11.0500	2
11.0700	1
11.1000	7
11.1800	6
11.2000	4
11.2500	2
11.3000	7
11.3300	5
11.3600	18
11.4000	4
11.4900	1
11.5000	31
11.5600	1
11.5700	6
11.6200	1
11.6800	4
11.7000	16
11.7200	2
11.7400	10
11.7900	4
11.8000	6
11.9000	3
12.0000	11
12.0500	27
12.0700	18
12.1000	2
12.1900	1
12.3000	4
12.3500	1
12.3900	5
12.4000	8
12.5000	3
12.5100	2
12.5500	26
12.6000	5
12.6200	3
12.6300	2
12.6500	2
12.7000	16
12.8000	49
12.8100	2
12.8300	2

12.8500	1
12.9000	17
12.9500	1
12.9700	1
12.9800	1
12.9900	46
13.0000	52
13.0100	7
13.0600	1
13.1000	14
13.1400	3
13.1700	1
13.2000	15
13.2200	4
13.2400	4
13.2900	1
13.3300	2
13.4000	10
13.4400	1
13.4900	2
13.5000	31
13.5300	2
13.6000	13
13.6800	15
13.7000	22
13.7300	3
13.8000	15
13.9000	4
13.9300	6
14.0000	20
14.0200	2
14.0700	1
14.1000	2
14.1600	14
14.2000	4
14.2100	12
14.2400	6
14.2800	16
14.3000	8
14.3300	1
14.3900	1
14.4000	10
14.4200	1
14.4700	1
14.4900	5
14.5300	15
14.5900	1
14.6000	2

14.6200	1
14.6600	4
14.6700	6
14.6900	6
14.7000	9
14.7400	11
14.7500	4
14.8000	13
14.8100	1
14.8400	17
14.9400	4
14.9500	4
15.0000	29
15.0400	15
15.0600	2
15.1000	46
15.1100	1
15.1500	2
15.1700	3
15.2000	5
15.2600	15
15.2900	14
15.3000	14
15.4000	15
15.4100	1
15.5000	14
15.6000	18
15.6300	4
15.6400	5
15.6800	4
15.7000	5
15.7300	15
15.7400	3
15.8000	25
15.8500	1
15.8700	2
15.9000	5
15.9600	7
15.9700	6
16.0000	63
16.0200	12
16.0500	1
16.0700	13
16.0900	7
16.1000	22
16.1200	1
16.2000	18
16.2500	2

16.3000	4
16.3600	12
16.3800	1
16.4000	5
16.4600	2
16.4700	41
16.5000	10
16.5100	1
16.5200	4
16.5500	23
16.6000	7
16.7300	6
16.7700	5
16.7800	7
16.8000	42
16.8200	4
16.9000	6
16.9300	2
16.9500	24
16.9600	5
16.9800	2
17.0000	126
17.0100	29
17.0500	15
17.0900	1
17.1000	20
17.1100	21
17.1600	1
17.1900	4
17.2000	7
17.2100	12
17.2400	1
17.3000	24
17.3200	3
17.4000	27
17.4300	2
17.4400	2
17.4500	2
17.5000	31
17.5400	2
17.5560	1
17.5700	14
17.6000	13
17.6700	3
17.6800	22
17.7000	8
17.7100	6
17.7200	2

17.8000	62
17.8400	1
17.8500	2
17.8800	3
17.9000	34
17.9200	14
17.9700	2
18.0000	53
18.0600	5
18.1000	7
18.1200	9
18.1500	4
18.1600	14
18.1900	3
18.2000	16
18.2300	1
18.2500	4
18.3000	5
18.3300	3
18.4000	5
18.4400	3
18.4800	5
18.4900	6
18.5000	43
18.5100	1
18.5300	3
18.5600	4
18.5900	2
18.6000	77
18.7000	24
18.7800	1
18.8000	2
18.8600	1
18.8800	20
18.9000	133
19.0000	30
19.0100	29
19.0800	9
19.0900	3
19.1000	27
19.1200	6
19.1500	3
19.1600	3
19.2000	7
19.2700	23
19.3000	27
19.3249	1
19.3300	5

19.3400	2
19.4000	31
19.4400	1
19.4900	1
19.5000	13
19.5900	3
19.6000	9
19.6400	8
19.6700	24
19.6900	2
19.7000	33
19.7100	3
19.7200	1
19.8100	33
19.8300	1
19.8700	14
19.9100	1
20.0000	48
20.0830	5
20.1400	19
20.3000	5
20.3400	8
20.3600	64
20.3700	3
20.3800	3
20.4000	21
20.4500	13
20.4600	4
20.5000	12
20.5100	19
20.5400	27
20.5800	3
20.6300	1
20.6400	4
20.6500	1
20.6800	6
20.7000	6
20.7300	17
20.7700	28
20.8500	7
20.8600	1
20.8900	3
20.9200	13
21.0000	1
21.0200	3
21.0300	2
21.1000	65
21.1200	5

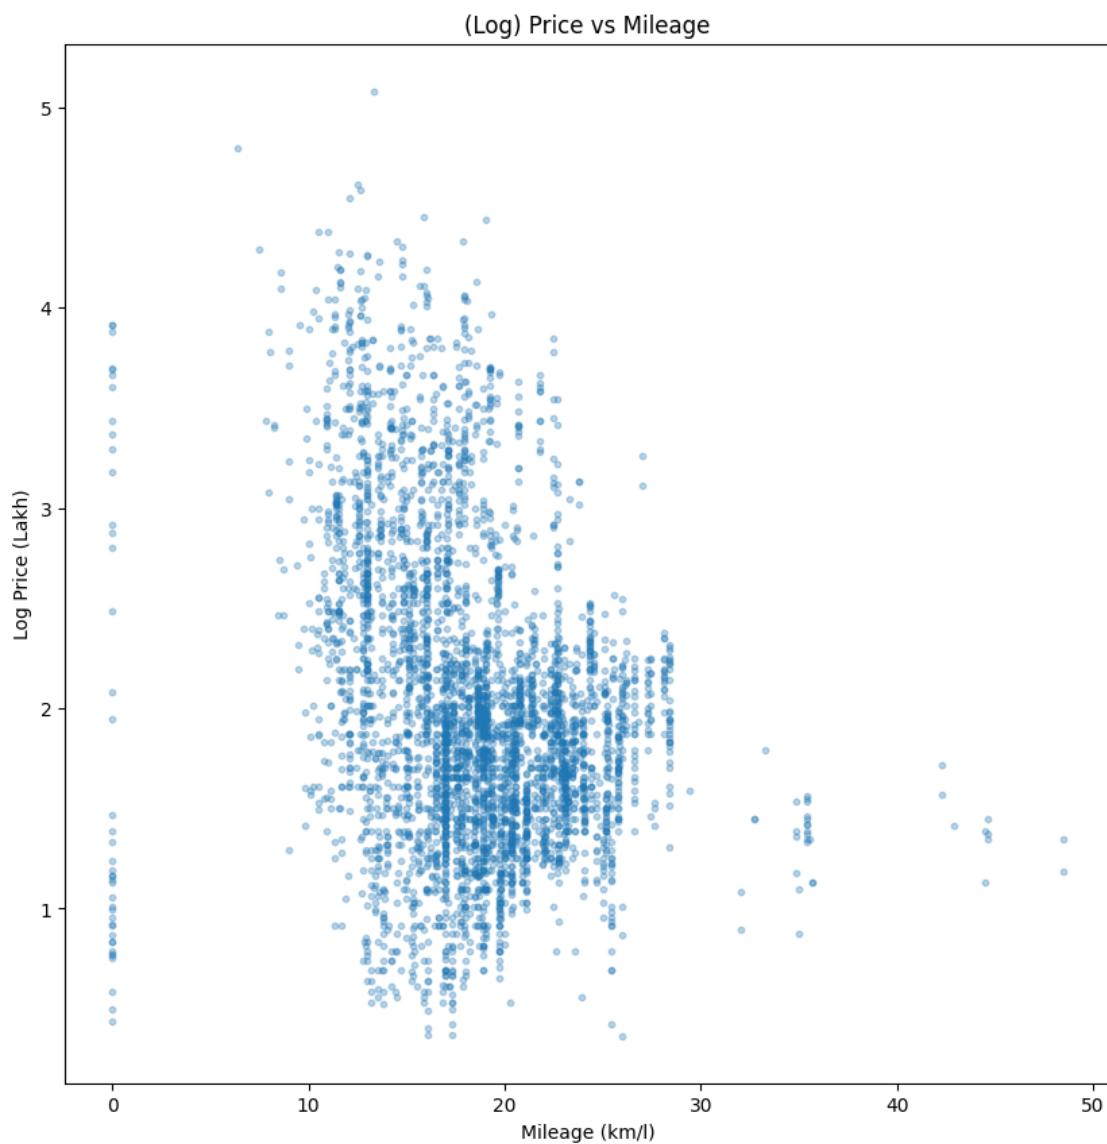
21.1300	1
21.1400	3
21.1900	4
21.2000	1
21.2100	3
21.2700	6
21.3800	2
21.4000	29
21.4300	9
21.5000	17
21.5600	2
21.6400	11
21.6600	5
21.7000	4
21.7600	10
21.7900	1
21.8000	1
21.9000	27
22.0000	18
22.0700	26
22.1000	2
22.3000	7
22.3200	39
22.4800	11
22.5000	19
22.5400	25
22.6100	1
22.6900	19
22.7000	49
22.7400	21
22.7700	10
22.9000	36
22.9500	6
23.0000	22
23.0100	1
23.0800	12
23.1000	39
23.2000	17
23.2750	2
23.3000	2
23.4000	28
23.5000	4
23.5700	1
23.5900	27
23.6500	5
23.8000	3
23.8400	7
23.9000	7

24.0000	35
24.0400	1
24.0700	23
24.2000	3
24.3000	35
24.4000	16
24.5000	3
24.5200	8
24.7000	14
24.8000	2
24.8825	1
25.0000	11
25.1000	14
25.1700	25
25.2000	26
25.3200	4
25.4000	13
25.4400	13
25.4700	7
25.6000	2
25.8000	39
25.8300	8
26.0000	25
26.1000	4
26.2100	9
26.5900	23
26.6000	1
26.8000	1
27.0300	2
27.2800	1
27.3000	13
27.3900	8
27.4000	2
27.6200	3
28.0900	17
28.4000	32
29.3930	1
32.0050	2
32.7180	2
33.2500	1
34.8460	4
34.9790	2
35.3780	11
35.5200	1
35.6839	2
42.2807	2
42.9058	1
44.4752	2

```
44.6082      3  
48.4700      2  
Name: count, dtype: int64
```

It is bell shaped as expected. However, the regular peaks hints that there was some bias towards round numbers.

```
[52]: plt.figure(figsize=(10, 10))  
plt.scatter(df_train.Mileage, np.log1p(df_train.Price), alpha=0.3, s=10)  
plt.xlabel("Mileage (km/l)")  
plt.ylabel("Log Price (Lakh)")  
plt.title("(Log) Price vs Mileage")  
plt.show()
```



Mileage alone does not show clear relationship with Price.

Cleaning up unrealistic outliers would stabilize the model.

```
[53]: class MileageClipper(TransformerMixin, BaseEstimator):
    def __init__(self,
                 mileage_col: str = "Mileage",
                 clipping_quantile: float = 0.995,
                 ):
        self.mileage_col = mileage_col
        self.clipping_quantile = clipping_quantile

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series] = None) -> "MileageClipper":
        return self

    def transform(self, X: pd.DataFrame):
        X = X.copy()
        X[self.mileage_col] = X[self.mileage_col].clip(
            upper=int(X[self.mileage_col].quantile(self.clipping_quantile)))
        )
        return X
```

```
[54]: mileage_outlier_clip_exp = Experiment(
    ExperimentConfig(
        name="clipping-extreme-mileage",
        pipeline=Pipeline(
            [
                ("mileage_clip_outliers", MileageClipper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

mileage_outlier_clip_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```
[Experiment: clipping-extreme-mileage]
Cross-validating (5-folds)...
CV score: 0.1578 ± 0.0153
-0.0028 -0.0010 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/clipping-extreme-mileage.csv
Experiment complete
```

```
[55]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("group_infrequent_fuel_type", FuelTypeGrouper()),
                ("mileage_clip_outliers", MileageClipper()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

current_best_exp.run(X_train, y_train, None, baseline_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1484 ± 0.0143
-0.0123 -0.0020 compared to baseline (Negative is better)
Training on full training set...
Experiment complete
```

## Engine

```
[56]: df_train.Engine.describe()
```

```
[56]: count      4428.0
mean       1618.2771
std        595.245047
min         624.0
25%        1198.0
50%        1493.0
75%        1968.0
max        5998.0
Name: Engine, dtype: float64
```

Engine feature seems reasonable enough.

## Power

```
[57]: df_train.Power.describe()
```

```
[57]: count    4336.000000
      mean     113.171817
      std      53.875993
      min     34.200000
      25%    75.000000
      50%    94.000000
      75%   138.100000
      max    560.000000
Name: Power, dtype: float64
```

**Power** feature seems reasonable enough.

However, as explored in `preprocessing.ipynb`, the percentage of missing values in `Power` was non-negligible.

```
[58]: df_train[df_train.Power.isna()]
```

	Name	Location	Year	Kilometers_Driven	Fuel_Type		
ID							
2CM572	Fiat Petra	Pune	2005	120000	Petrol		
4J1SFY	Mercedes-Benz E-Class	Pune	2001	121000	Diesel		
LHXSLV	Maruti Swift	Hyderabad	2014	81609	Diesel		
R7UPR3	Fiat Siena	Jaipur	2001	70000	Petrol		
KZI5XI	Skoda Laura	Pune	2010	85000	Petrol		
...	...	...	...	...	...		
Z3GW3N	Hyundai Santro	Hyderabad	2006	74483	Petrol		
M49PV9	Hyundai Santro	Mumbai	2005	102000	Petrol		
9BPQA2	Hyundai Santro	Pune	2005	100000	CNG		
AOLNGH	Hyundai Santro	Jaipur	2004	200000	Petrol		
MDPHOT	Hyundai Santro	Kochi	2007	58815	Petrol		
	Transmission	Owner_Type	Mileage	Engine	Power		
ID							
2CM572	Manual	Second	15.50	1242	Nan	Others	5
4J1SFY	Manual	First	15.00	2148	Nan	White	5
LHXSLV	Manual	First	17.80	1248	Nan	Others	5
R7UPR3	Manual	Third	0.00	1242	Nan	White	5
KZI5XI	Manual	First	17.50	1798	Nan	Black/Silver	5
...	...	...	...	...	...	...	
Z3GW3N	Automatic	First	0.00	999	Nan	Black/Silver	5
M49PV9	Manual	Second	17.00	1086	Nan	White	5
9BPQA2	Manual	Third	22.61	1086	Nan	White	5
AOLNGH	Manual	First	0.00	1086	Nan	Black/Silver	5
MDPHOT	Manual	First	17.00	1086	Nan	White	5
	Doors	New_Price	Price				
ID							
2CM572	4	Nan	0.85				

```

4J1SFY      4      NaN  5.00
LHXSLV      4      NaN  5.55
R7UPR3      4      NaN  0.55
KZI5XI      4      NaN  2.85
...
...       ...    ...
Z3GW3N      4      NaN  2.30
M49PV9      4      NaN  0.85
9BPQA2      4      NaN  1.20
AOLNGH      4      NaN  0.80
MDPHOT      4      NaN  1.99

```

[92 rows x 15 columns]

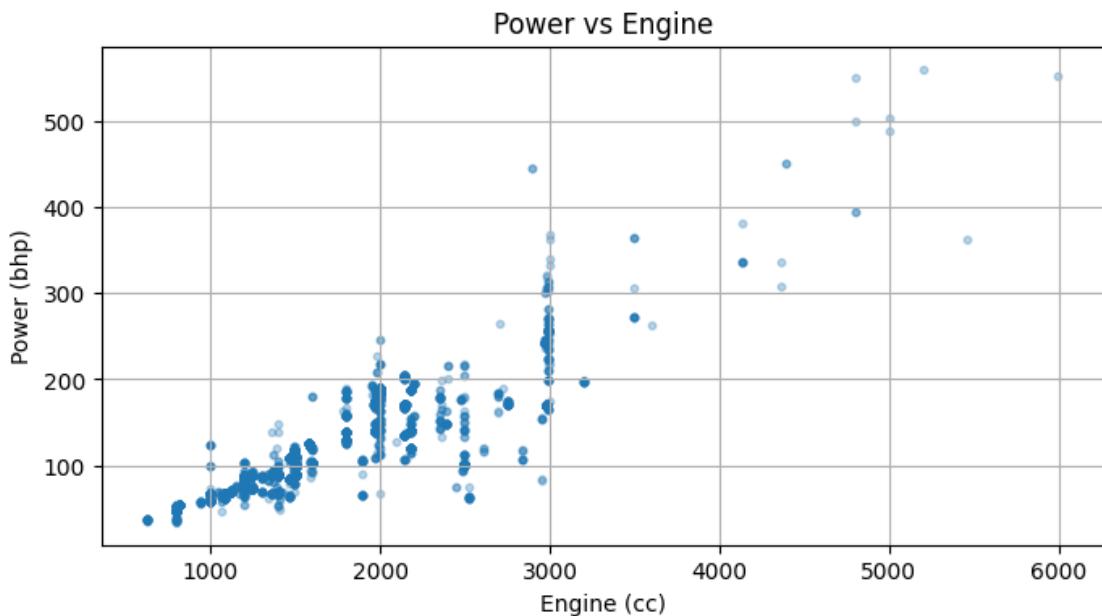
Power can be faithfully imputed exploiting its high relevance to Engine.

```
[59]: print(f"Engine-Power Correlation: {df_train.Engine.corr(df_train.Power):.2f}")
```

Engine-Power Correlation: 0.86

```
[60]: plt.figure(figsize=(8, 4))
plt.scatter(df_train["Engine"], df_train["Power"], alpha=0.3, s=10)

plt.xlabel("Engine (cc)")
plt.ylabel("Power (bhp)")
plt.title("Power vs Engine")
plt.grid(True)
plt.show()
```



```
[61]: class PowerImputer(TransformerMixin, BaseEstimator):
    def __init__(self, engine_col: str = "Engine", power_col="Power", clip_negative: bool = True):
        self.engine_col = engine_col
        self.power_col = power_col
        self.clip_negative = clip_negative

    def fit(self, X: pd.DataFrame, y=Optional[pd.Series]) -> "PowerImputer":
        df = X[[self.engine_col, self.power_col]].dropna()
        engines = df[self.engine_col].astype(float)
        powers = df[self.power_col].astype(float)

        # Linear regression
        cov = ((engines - engines.mean()) * (powers - powers.mean())).sum()
        var = ((engines - engines.mean()) ** 2).sum()
        self.slope = cov / var
        self.intercept = powers.mean() - self.slope * engines.mean()

        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()

        mask = X[self.power_col].isna()
        has_engine = X[self.engine_col].notna() & mask
        X.loc[has_engine, self.power_col] = (
            self.slope * X.loc[has_engine, self.engine_col] + self.intercept
        )

        if self.clip_negative:
            X[self.power_col] = X[self.power_col].clip(lower=0)

        return X
```

```
[62]: power_imputation_exp = Experiment(
    ExperimentConfig(
        name="power-imputation",
        pipeline=Pipeline(
            [
                ("imput_power", PowerImputer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)
```

```
)
```

```
power_imputation_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```
[Experiment: power-imputation]
Cross-validating (5-folds)...
CV score: 0.1551 ± 0.0159
    -0.0055 -0.0004 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/power-imputation.csv
Experiment complete
```

```
[63]: imputer = PowerImputer()
imputer.fit(X_train)

slope = imputer.slope
intercept = imputer.intercept

df_plot = X_train.copy()
df_plot["imputed_Power"] = imputer.transform(X_train)[imputer.power_col]
df_plot["missing_Power"] = df_plot.Power.isna()

engine_range = np.linspace(df_plot.Engine.min(), df_plot.Engine.max(), 200)
regression_line = slope * engine_range + intercept

observed_mask = ~df_plot.missing_Power

df_observed = df_plot.loc[observed_mask]
residuals = df_observed.Power - (slope * df_observed.Engine + intercept)
sigma = residuals.std()

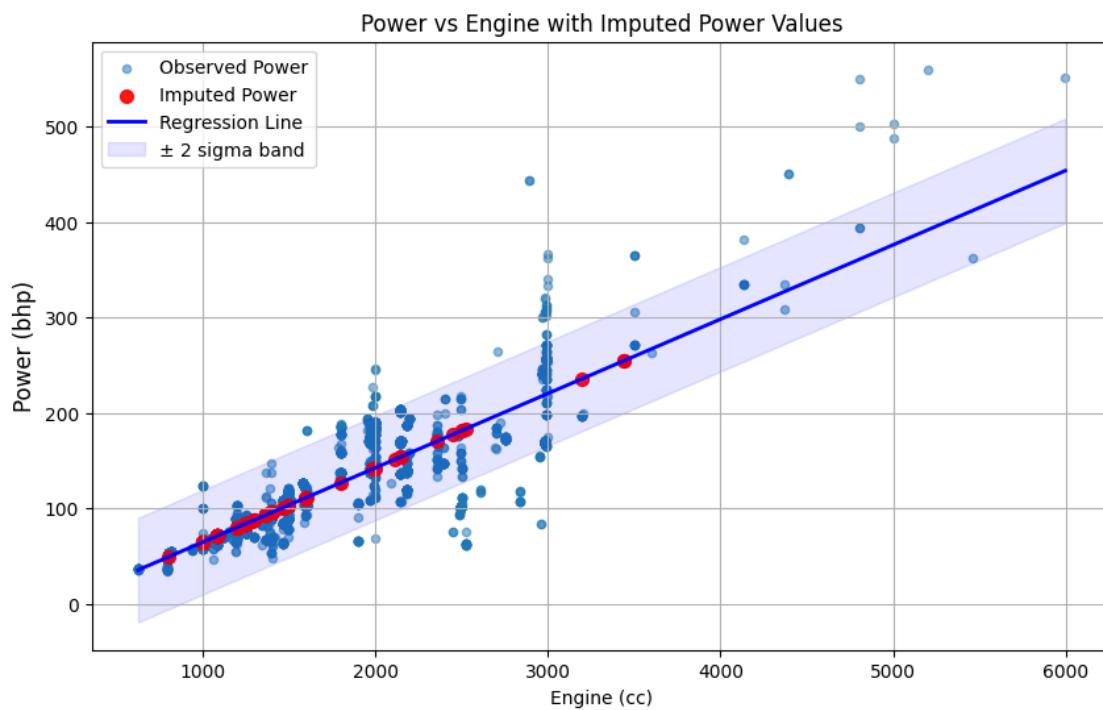
plt.figure(figsize=(10, 6))
plt.scatter(
    df_plot.loc[observed_mask, "Engine"],
    df_plot.loc[observed_mask, "Power"],
    s=20,
    alpha=0.5,
    label="Observed Power",
)
plt.scatter(
    df_plot.loc[df_plot.missing_Power, "Engine"],
    df_plot.loc[df_plot.missing_Power, "imputed_Power"],
    s=50,
    color="red",
    alpha=0.9,
```

```

        label="Imputed Power",
)
plt.plot(
    engine_range, regression_line, color="blue", linewidth=2, label="Regression Line"
)
plt.fill_between(
    engine_range,
    regression_line - 2 * sigma,
    regression_line + 2 * sigma,
    color="blue",
    alpha=0.1,
    label="± 2 sigma band",
)

plt.xlabel("Engine (cc)", fontsize=10)
plt.ylabel("Power (bhp)", fontsize=12)
plt.title("Power vs Engine with Imputed Power Values")
plt.legend()
plt.grid(True)
plt.show()

```



Missing **Power** was effectively imputed with its highly linear correlation with **Engine**.

Combining all feature engineering so far.

```
[64]: df_train.columns
```

```
[64]: Index(['Name', 'Location', 'Year', 'Kilometers_Driven', 'Fuel_Type',
       'Transmission', 'Owner_Type', 'Mileage', 'Engine', 'Power', 'Colour',
       'Seats', 'Doors', 'New_Price', 'Price'],
      dtype='object')
```

```
[65]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("group_infrequent_fuel_type", FuelTypeGrouper()),
                ("mileage_clip_outliers", MileageClipper()),
                ("imput_power", PowerImputer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    ),
)

current_best_exp.run(X_train, y_train, None, baseline_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1466 ± 0.0165
-0.0140 +0.0002 compared to baseline (Negative is better)
Training on full training set...
Experiment complete
```

## Colour

```
[66]: df_train.Colour.describe()
```

```
[66]: count      4428
unique         3
top      White
freq      1539
Name: Colour, dtype: object
```

```
[67]: df_train.Colour.value_counts()
```

```
[67]: Colour
White           1539
Others          1519
Black/Silver    1370
Name: count, dtype: int64
```

Colour is already well balanced.

### 0.0.3 Seats

```
[68]: df_train.Seats.value_counts()
```

```
[68]: Seats
5      3736
7      503
8      83
4      69
6      21
2      12
10     2
9      2
Name: count, dtype: Int64
```

Binning could help for those rare seat counts. However, the binning should not break the ordinality.

```
[69]: class SeatsBinner(TransformerMixin, BaseEstimator):
    def __init__(self, seats_col: str = "Seats"):
        self.seats_col = seats_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> "SeatsBinner":
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()
        seats = X[self.seats_col]

        X[self.seats_col] = seats.replace(
            {
                2: 4,
                4: 4,
                # Small
                5: 5,
                # Standard
                6: 7,
                7: 7,
                8: 7,
```

```

        # Large
        9: 9,
        10: 9,
        # Van
    }
)

return X

```

```
[70]: seats_binning_exp = Experiment(
    ExperimentConfig(
        name="bin-seats",
        pipeline=Pipeline(
            [
                ("bin_seats", SeatsBinner()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

seats_binning_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

[Experiment: bin-seats]  
Cross-validating (5-folds)...  
CV score: 0.1581 ± 0.0154  
-0.0025 -0.0009 compared to baseline (Negative is better)  
Training on full training set...  
Creating submission on test set...  
Submission created: artifacts/experiment-results/bin-seats.csv  
Experiment complete

Binning Seats had positive effect on the prediction score.

```
[71]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
            ]
        )
    )
)
```

```

        ("group_infrequent_fuel_type", FuelTypeGrouper()),
        ("mileage_clip_outliers", MileageClipper()),
        ("imput_power", PowerImputer()),
        ("bin_seats", SeatsBinner()),
        ("category_encode", CategoricalEncoder()),
        ("model", XGBRegressor()),
    ]
),
)
)

current_best_exp.run(X_train, y_train, None, baseline_exp_result);

```

[Experiment: current-best]  
Cross-validating (5-folds)...  
CV score: 0.1479 ± 0.0159  
-0.0127 -0.0004 compared to baseline (Negative is better)  
Training on full training set...  
Experiment complete

## Doors

[72]: df\_train.Doors.value\_counts()

[72]: Doors

4	3884
5	532
2	12
Name:	count, dtype: Int64

There are very small number of cars with 2 doors. However, I think it should not be binned, since usually cars with 2 doors are very expensive sports car, and it can imply a lot, even though it has very small sample size.

[73]: df\_train[df\_train.Doors == 2][["Name", "Engine", "Power", "Seats", "Doors"]]

[73]:

ID	Name	Engine	Power	Seats	Doors
UFNCV8	BMW Z4	2979	306.00	2	2
ICAVC1	Jaguar F	5000	488.10	2	2
YDPHR8	Smart Fortwo	799	NaN	2	2
RS5FNO	Audi A4	3197	NaN	2	2
X5JJRY	Porsche Boxster	2706	265.00	2	2
20AGHR	Porsche Cayman	3436	NaN	2	2
ORNZ40	Porsche Cayman	3436	NaN	2	2
W15E5A	Mercedes-Benz SLK-Class	3498	306.00	2	2
NF6IF1	BMW Z4	2979	306.00	2	2
RGDCTT	Mercedes-Benz SLC	2996	362.07	2	2

IJ96ZM	Audi TT	1984	207.80	2	2
PYZN3W	Lamborghini Gallardo	5204	560.00	2	2

### New\_Price

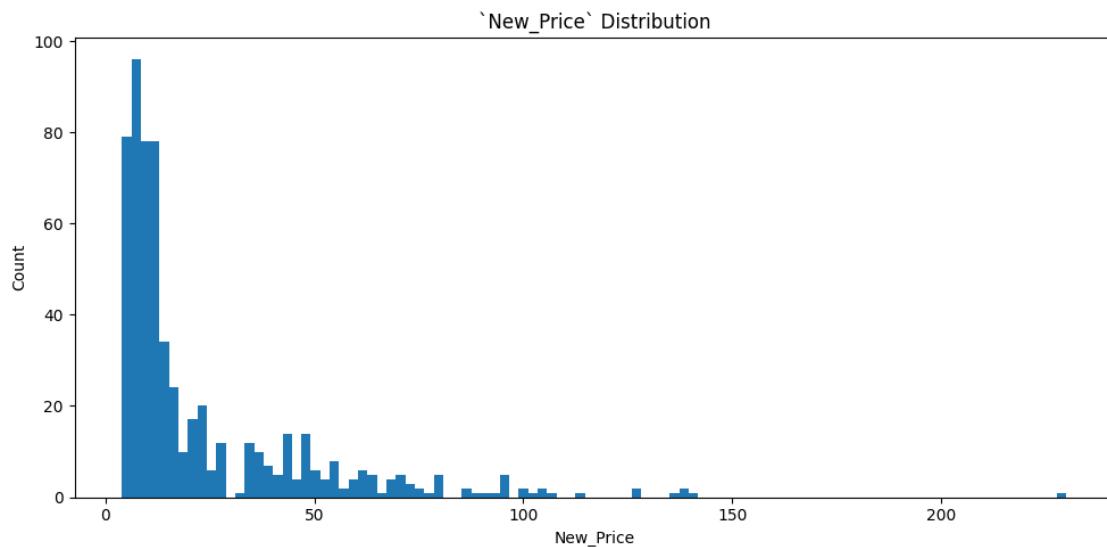
```
[74]: print(
    f"# of missing New_Price: {len(df_train[df_train.New_Price.isna()])} out of {len(df_train)}"
)
```

# of missing New\_Price: 3827 out of 4428

```
[75]: observed_new_prices = df_train[df_train.New_Price.notna()].New_Price

plt.figure(figsize=(10, 5))
plt.hist(observed_new_prices, bins=100)
plt.xlabel("New_Price")
plt.ylabel("Count")
plt.title("`New_Price` Distribution")

plt.tight_layout()
plt.show()
```



Due to its high sparsity, it would be dangerous to impute naively. Instead, labeling missing New\_Price might give non-trivial information from the pattern of missing New\_Price. Additionally, log transformation could help the heavy skew in its distribution.

```
[76]: observed_new_prices = df_train[df_train.New_Price.notna()].New_Price

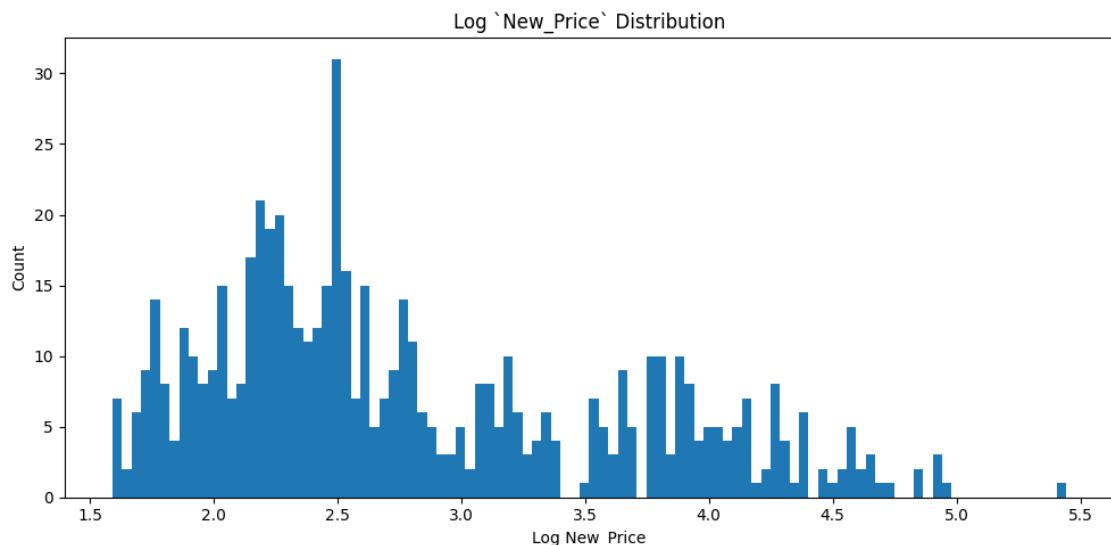
plt.figure(figsize=(10, 5))
```

```

plt.hist(np.log1p(observe_new_prices), bins=100)
plt.xlabel("Log New_Price")
plt.ylabel("Count")
plt.title("Log `New_Price` Distribution")

plt.tight_layout()
plt.show()

```



```

[77]: class NewPriceTransformer(TransformerMixin, BaseEstimator):
    def __init__(self, new_price_col: str = "New_Price"):
        self.new_price_col = new_price_col
        self.missing_col = "Missing_" + new_price_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> Union["NewPriceTransformer", None]:
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()

        X[self.missing_col] = X[self.new_price_col].isna().astype(int)

        X[self.new_price_col] = X[self.new_price_col].apply(np.log1p)

        return X

```

```
[78]: new_price_transform_exp = Experiment(
    ExperimentConfig(

```

```

        name="transform-new-price",
        pipeline=Pipeline(
            [
                ("transform_new_price", NewPriceTransformer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

new_price_transform_exp.run(X_train, y_train, X_test, baseline_exp_result);

```

```

[Experiment: transform-new-price]
Cross-validating (5-folds)...
CV score: 0.1606 ± 0.0163
    +0.0000  +0.0000 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/transform-new-price.csv
Experiment complete

```

It did not yield visible gain. However, I'll keep it since this transformation can be useful in making interaction features. Should be judged later.

```

[79]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                ("extract_brand_model", BrandModelExtractor()),
                (
                    "target_encode",
                    TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
                ),
                ("transform", YearToAgeTransformer()),
                ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
                ("group_infrequent_fuel_type", FuelTypeGrouper()),
                ("mileage_clip_outliers", MileageClipper()),
                ("imput_power", PowerImputer()),
                ("bin_seats", SeatsBinner()),
                ("transform_new_price", NewPriceTransformer()),
                ("category_encode", CategoricalEncoder()),
                ("model", XGBRegressor()),
            ]
        ),
    )
)

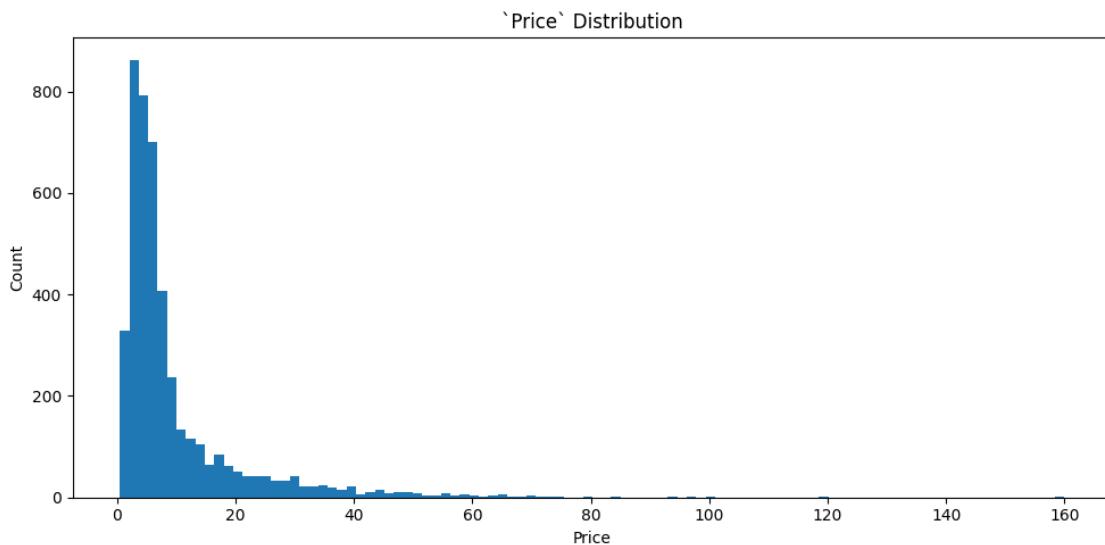
```

```
)  
  
current_best_exp.run(X_train, y_train, None, baseline_exp_result);
```

```
[Experiment: current-best]  
Cross-validating (5-folds)...  
CV score: 0.1479 ± 0.0159  
    -0.0127 -0.0004 compared to baseline (Negative is better)  
Training on full training set...  
Experiment complete
```

### Price

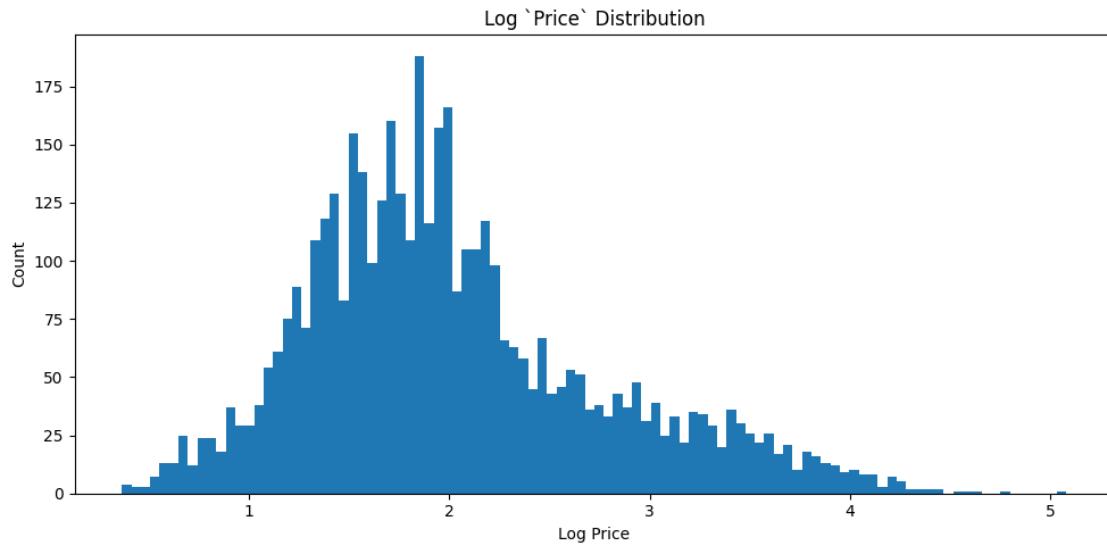
```
[80]: plt.figure(figsize=(10, 5))  
plt.hist(df_train.Price, bins=100)  
plt.xlabel("Price")  
plt.ylabel("Count")  
plt.title("`Price` Distribution")  
  
plt.tight_layout()  
plt.show()
```



The target feature `Price` is heavily skewed. Log transformation can mediate this.

```
[81]: plt.figure(figsize=(10, 5))  
plt.hist(np.log1p(df_train.Price), bins=100)  
plt.xlabel("Log Price")  
plt.ylabel("Count")  
plt.title("Log `Price` Distribution")
```

```
plt.tight_layout()  
plt.show()
```



```
[82]: price_transform_exp = Experiment(  
    ExperimentConfig(  
        name="transform-price",  
        pipeline=Pipeline(  
            [  
                ("category_encode", CategoricalEncoder()),  
                (  
                    "model",  
                    TransformedTargetRegressor(  
                        regressor=XGBRegressor(), func=np.log1p,  
                        inverse_func=np.expm1  
                    ),  
                    ),  
                ),  
            ],  
        ),  
    )  
)  
  
price_transform_exp.run(X_train, y_train, X_test, baseline_exp_result);
```

```
[Experiment: transform-price]  
Cross-validating (5-folds)...  
CV score: 0.1456 ± 0.0170  
-0.0150 +0.0007 compared to baseline (Negative is better)  
Training on full training set...
```

```
Creating submission on test set...
Submission created: artifacts/experiment-results/transform-price.csv
Experiment complete
```

Log transformation on Price was very effective.

#### 0.0.4 Combining best transformations for each features

```
[83]: best_single_feature_pipeline = Pipeline(
    [
        ("extract_brand_model", BrandModelExtractor()),
        (
            "target_encode",
            TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),
        ),
        ("transform", YearToAgeTransformer()),
        ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),
        ("group_infrequent_fuel_type", FuelTypeGrouper()),
        ("mileage_clip_outliers", MileageClipper()),
        ("imput_power", PowerImputer()),
        ("bin_seats", SeatsBinner()),
        ("transform_new_price", NewPriceTransformer()),
        ("category_encode", CategoricalEncoder()),
        (
            "model",
            TransformedTargetRegressor(
                regressor=XGBRegressor(), func=np.log1p, inverse_func=np.expm1
            ),
        ),
    ],
)
```

```
[84]: best_feature_by_feature_exp = Experiment(
    ExperimentConfig(
        name="combine-all-feature-by-feature-engineering",
        pipeline=best_single_feature_pipeline,
    )
)

best_feature_by_feature_exp_result = best_feature_by_feature_exp.run(
    X_train, y_train, X_test, baseline_exp_result
)
```

```
[Experiment: combine-all-feature-by-feature-engineering]
Cross-validating (5-folds)...
CV score: 0.1379 ± 0.0173
```

```

-0.0228 +0.0010 compared to baseline (Negative is better)
Training on full training set...
Creating submission on test set...
Submission created: artifacts/experiment-results/combine-all-feature-by-feature-
engineering.csv
Experiment complete

```

---

### 0.0.5 Interactions

To determine which interactions are worth investigating, I first perform residual diagnostics.

```
[85]: best_single_feature_pipeline.fit(X_train, y_train)

prediction_by_best_single_feature_pipeline = pd.Series(
    best_single_feature_pipeline.predict(X_train), index=y_train.index
)
residuals = y_train - prediction_by_best_single_feature_pipeline

X_train_transformed = pd.DataFrame(
    best_single_feature_pipeline[:-1].fit_transform(X_train, y_train)
)
y_train_transformed = y_train.apply(np.log10)

df_residuals = X_train_transformed.copy()
df_residuals["Residual"] = residuals
df_residuals["Prediction"] = prediction_by_best_single_feature_pipeline
df_residuals["Price"] = y_train

df_residuals
```

	Name	Location	Kilometers_Driven	Fuel_Type	Transmission	\			
ID									
G4XLU0	4.267524	15.550690	59138	0	1				
CRSHOS	7.206730	11.252232	81504	0	1				
FUJ4X1	6.443593	10.093132	92000	2	1				
QMVK6E	5.077432	5.460807	33249	0	1				
4SWHFC	5.867476	12.628270	65000	2	1				
...	...	...	...	...	...				
TR7SLB	10.542137	11.252232	51884	0	1				
QB41QE	6.497267	5.460807	27210	0	1				
ODG8N7	33.110594	6.749484	52000	0	0				
EV2ZBX	2.542287	10.096643	56000	2	1				
J2RCU8	18.250472	12.628270	52000	0	0				
	Owner_Type	Mileage	Engine	Power	Colour	Seats	Doors	New_Price	\

ID									
G4XLU0	0	17.00	1405	70.00	1	5	4	NaN	
CRSHOS	0	21.43	1364	87.20	1	5	4	NaN	
FUJ4X1	0	13.80	1299	70.00	1	5	4	NaN	
QMVK6E	0	21.27	1396	88.76	0	5	4	NaN	
4SWHFC	0	17.00	1497	118.00	2	5	4	NaN	
...	...	...	...	...	...	...	...	...	
TR7SLB	0	16.00	2179	140.00	2	7	5	NaN	
QB41QE	0	27.30	1498	98.60	1	5	4	NaN	
ODG8N7	0	12.70	2179	187.70	2	5	4	NaN	
EV2ZBX	0	24.70	796	47.30	1	5	4	NaN	
J2RCU8	0	12.00	2987	224.00	0	7	5	NaN	

	Brand	Model	Age	Missing_New_Price	Residual	Prediction	\
ID							
G4XLU0	3.417323	4.267524	7		1 0.021814	2.558186	
CRSHOS	11.417767	7.206730	7		1 -0.088608	6.618608	
FUJ4X1	6.549908	6.443593	13		1 -0.100111	1.350111	
QMVK6E	5.333501	5.077432	8		1 -0.016794	3.266794	
4SWHFC	5.425303	5.867476	9		1 0.109849	5.090151	
...	...	...	...	...	...	...	
TR7SLB	8.075126	10.542137	4		1 0.440084	12.019916	
QB41QE	5.425303	6.497267	4		1 -0.017071	5.867071	
ODG8N7	37.771935	33.110594	5		1 0.467415	39.282585	
EV2ZBX	4.604672	2.542287	7		1 -0.037402	2.137402	
J2RCU8	27.104100	18.250472	6		1 2.699703	46.300297	

	Price
ID	
G4XLU0	2.58
CRSHOS	6.53
FUJ4X1	1.25
QMVK6E	3.25
4SWHFC	5.20
...	...
TR7SLB	12.46
QB41QE	5.85
ODG8N7	39.75
EV2ZBX	2.10
J2RCU8	49.00

[4428 rows x 20 columns]

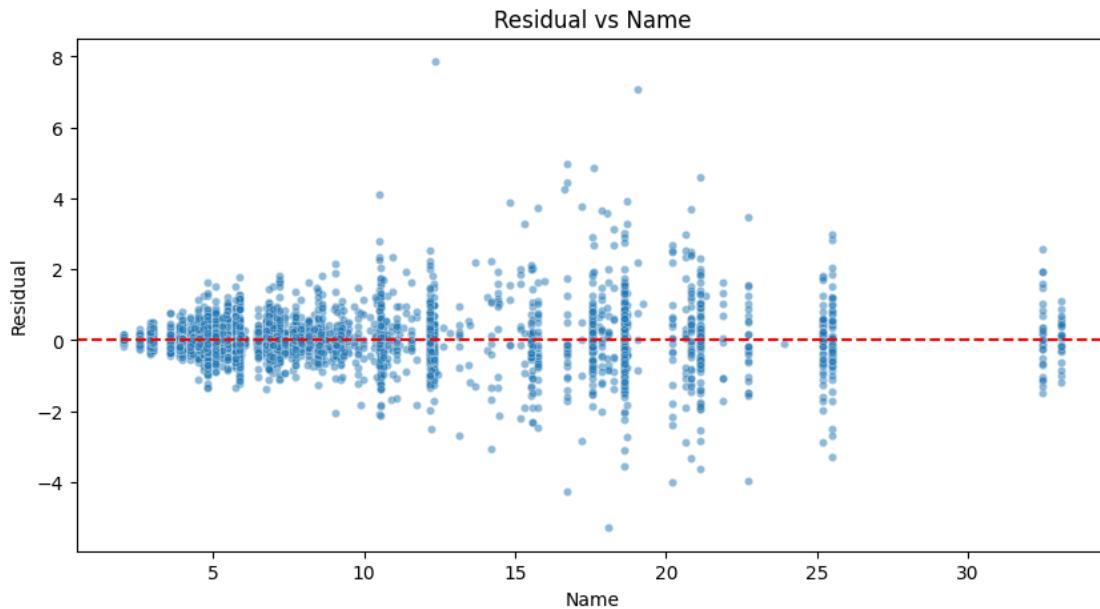
```
[86]: def residual_plot(feature, bins=False):
    plt.figure(figsize=(10, 10 if bins else 5))
    x = df_residuals[feature][df_residuals[feature].notna()]
    if bins:
```

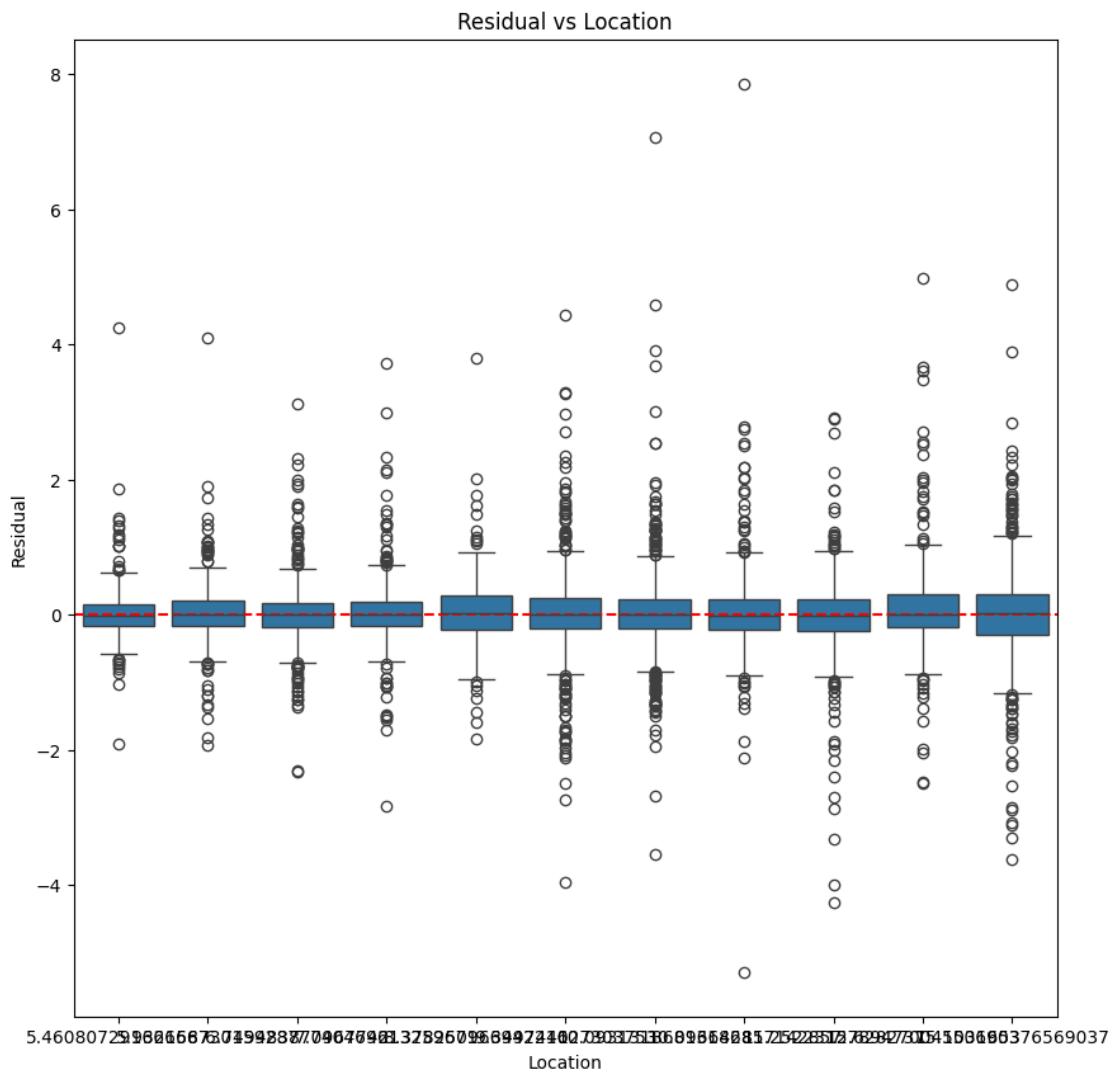
```

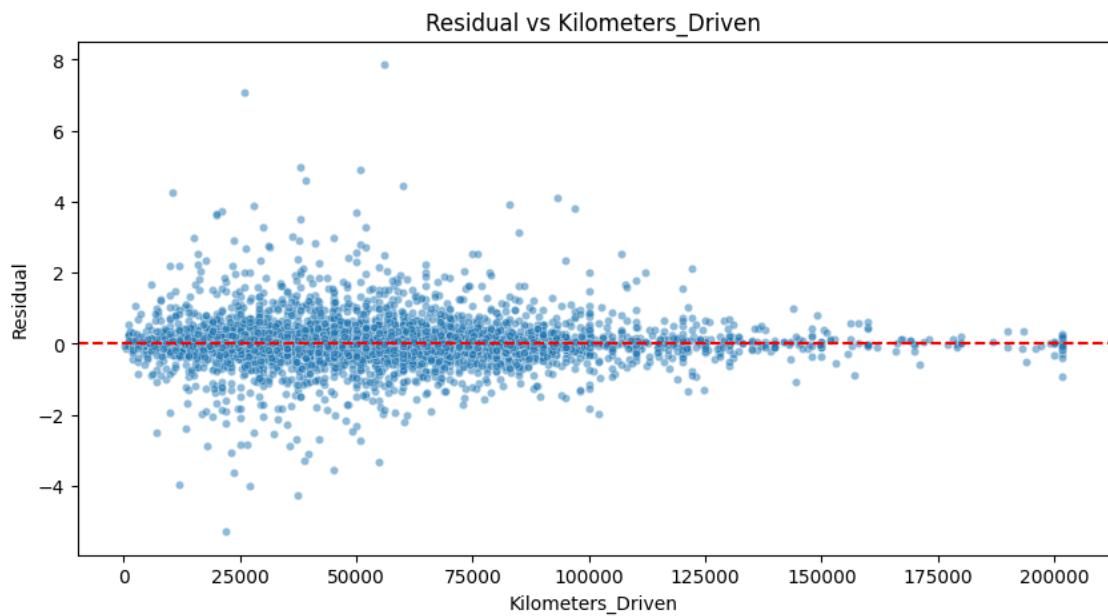
        sns.boxplot(x=x, y=df_residuals["Residual"])
    else:
        sns.scatterplot(
            x=x,
            y=df_residuals["Residual"],
            alpha=0.5,
            s=20,
        )
    plt.axhline(0, color="red", linestyle="--")
    plt.title(f"Residual vs {feature}")
    plt.show()

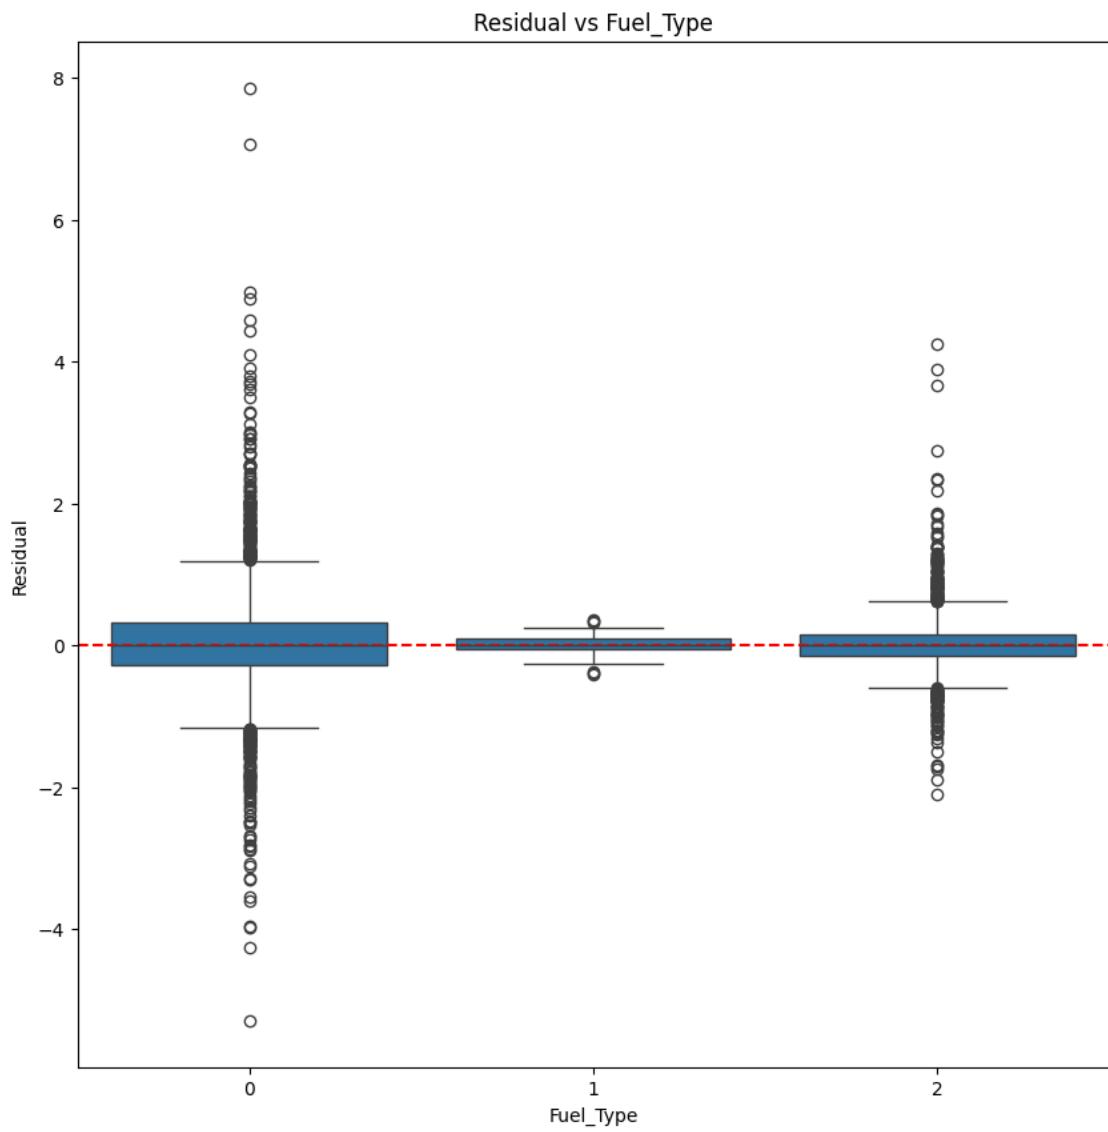
```

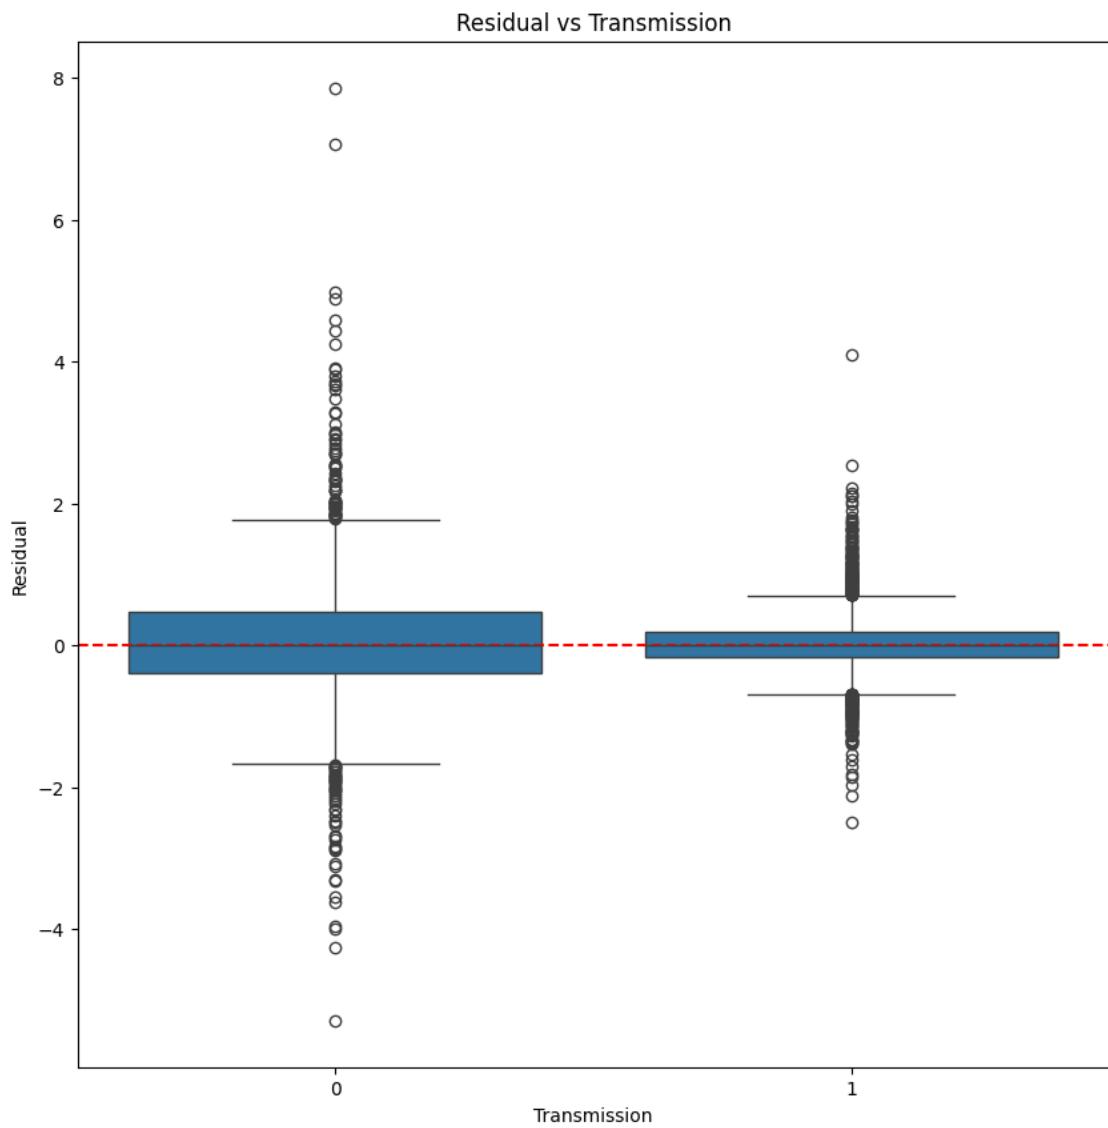
```
[87]: for feature in df_residuals.drop(columns=["Residual", "Price", "Prediction"]):
    columns:
        use_box_plot = (
            True
            if len(df_residuals[feature].unique()) < 15
            or df_residuals[feature].dtype.name == "category"
            else False
        )
    residual_plot(feature, bins=use_box_plot)
```



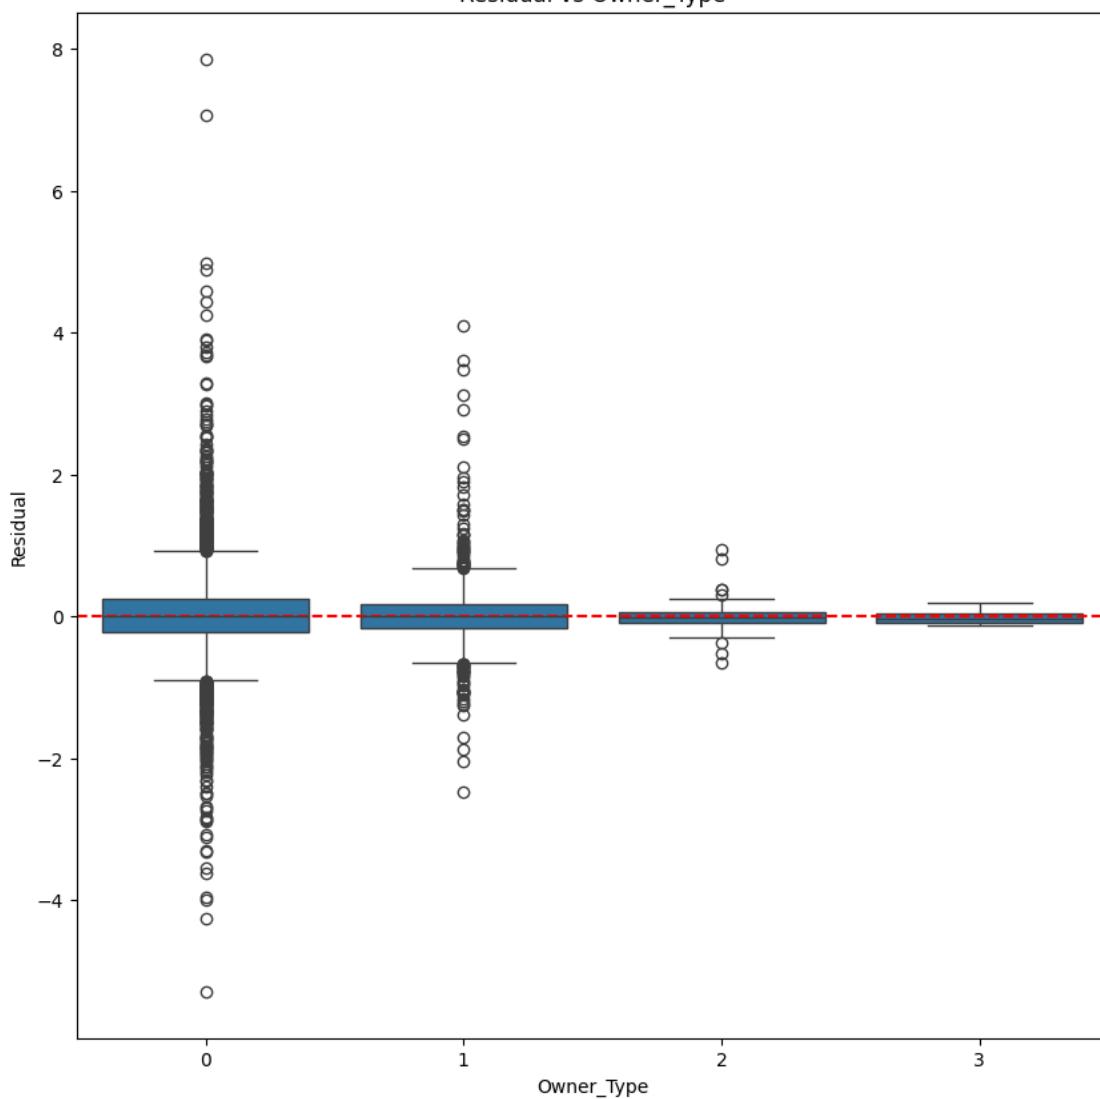




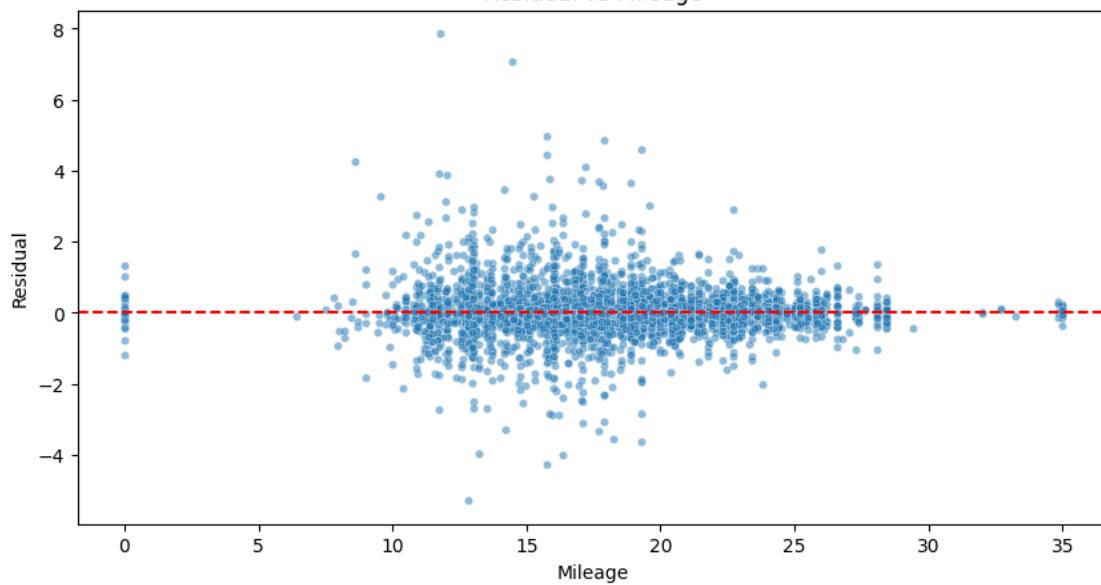




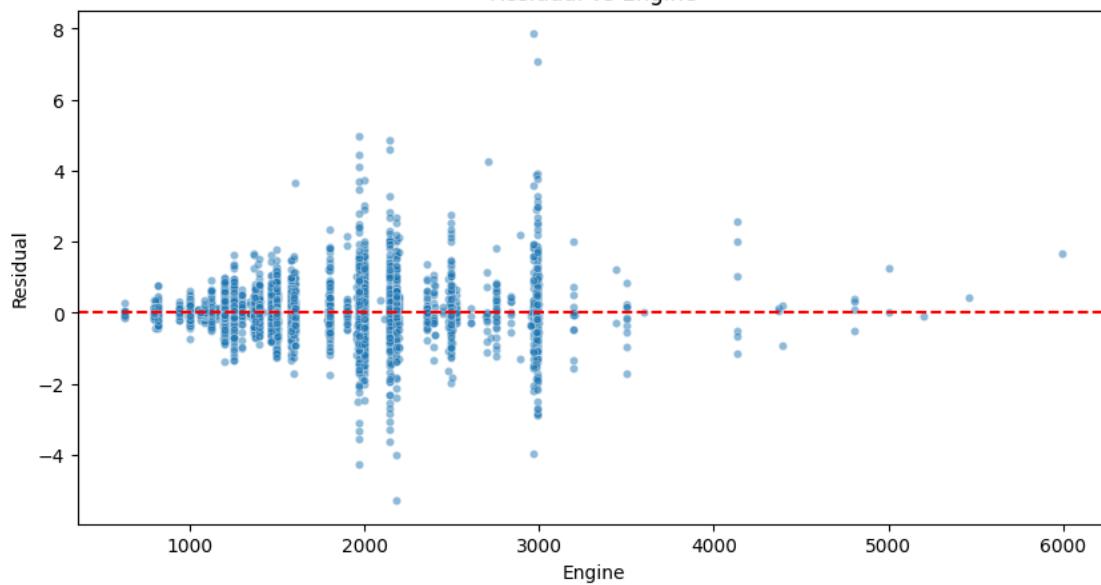
Residual vs Owner\_Type

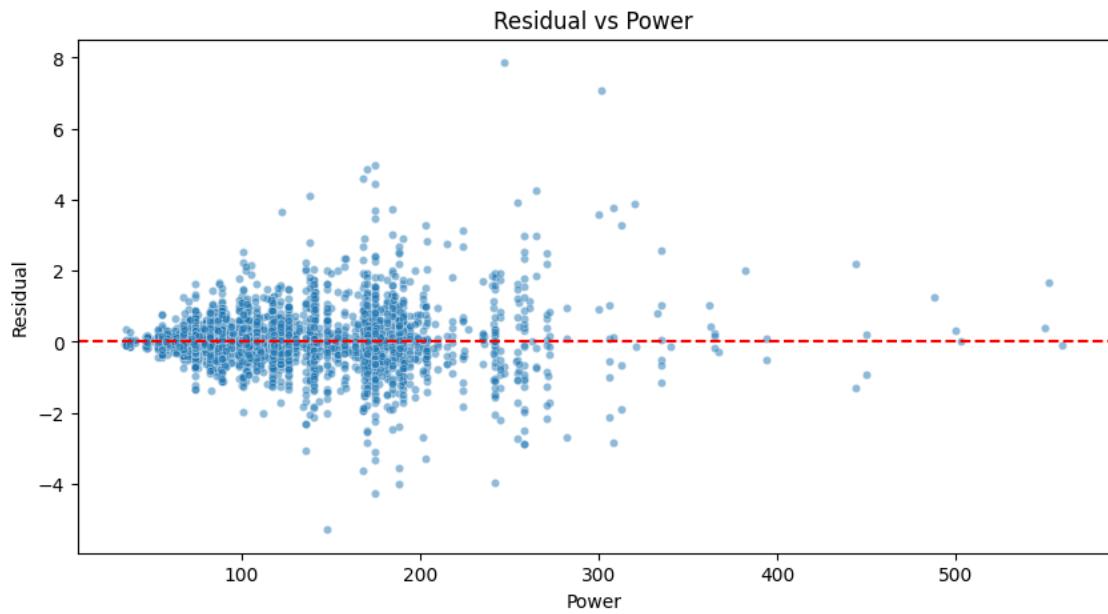


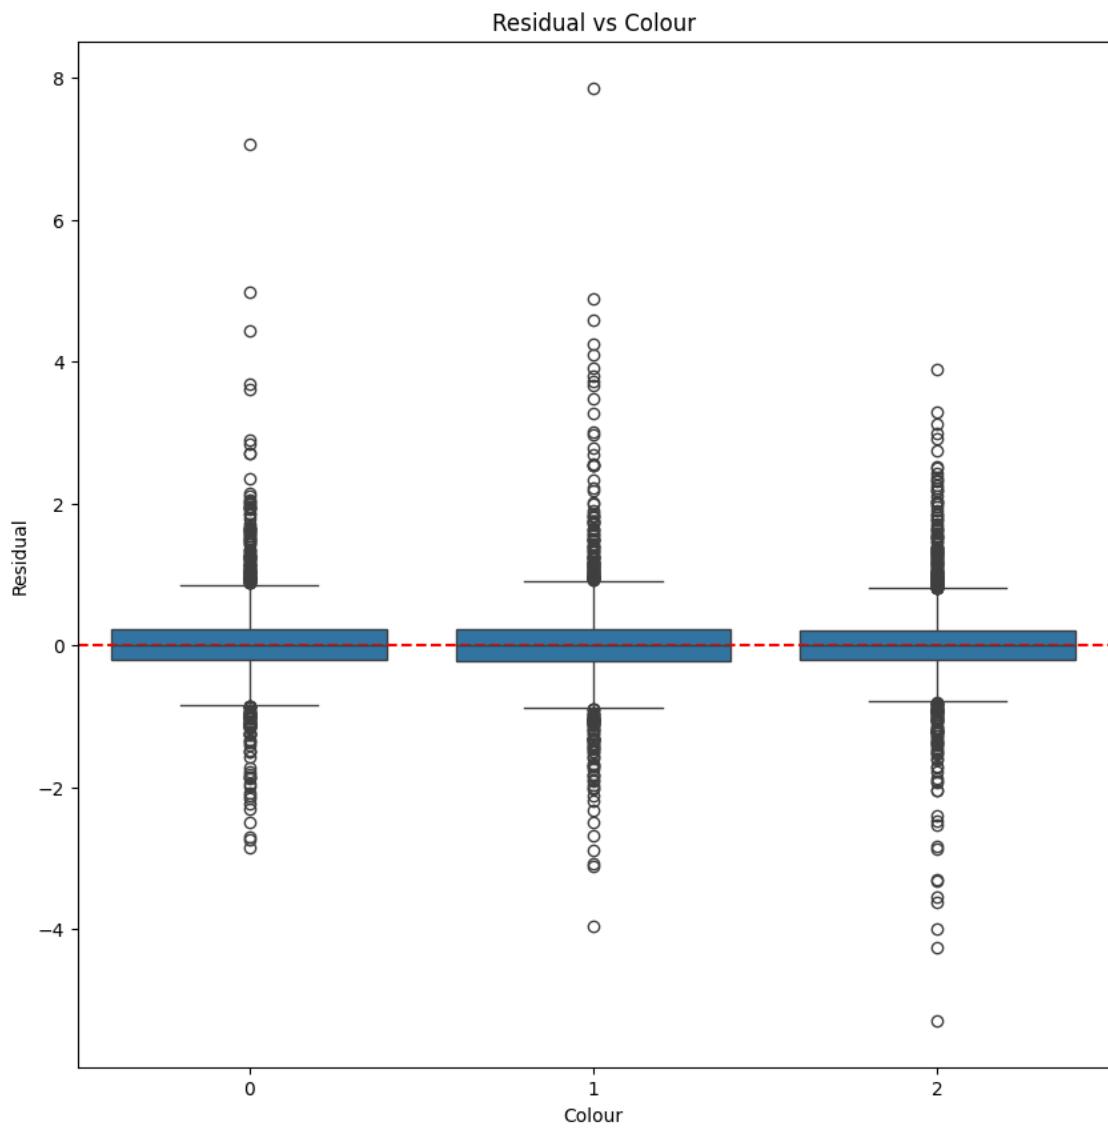
Residual vs Mileage

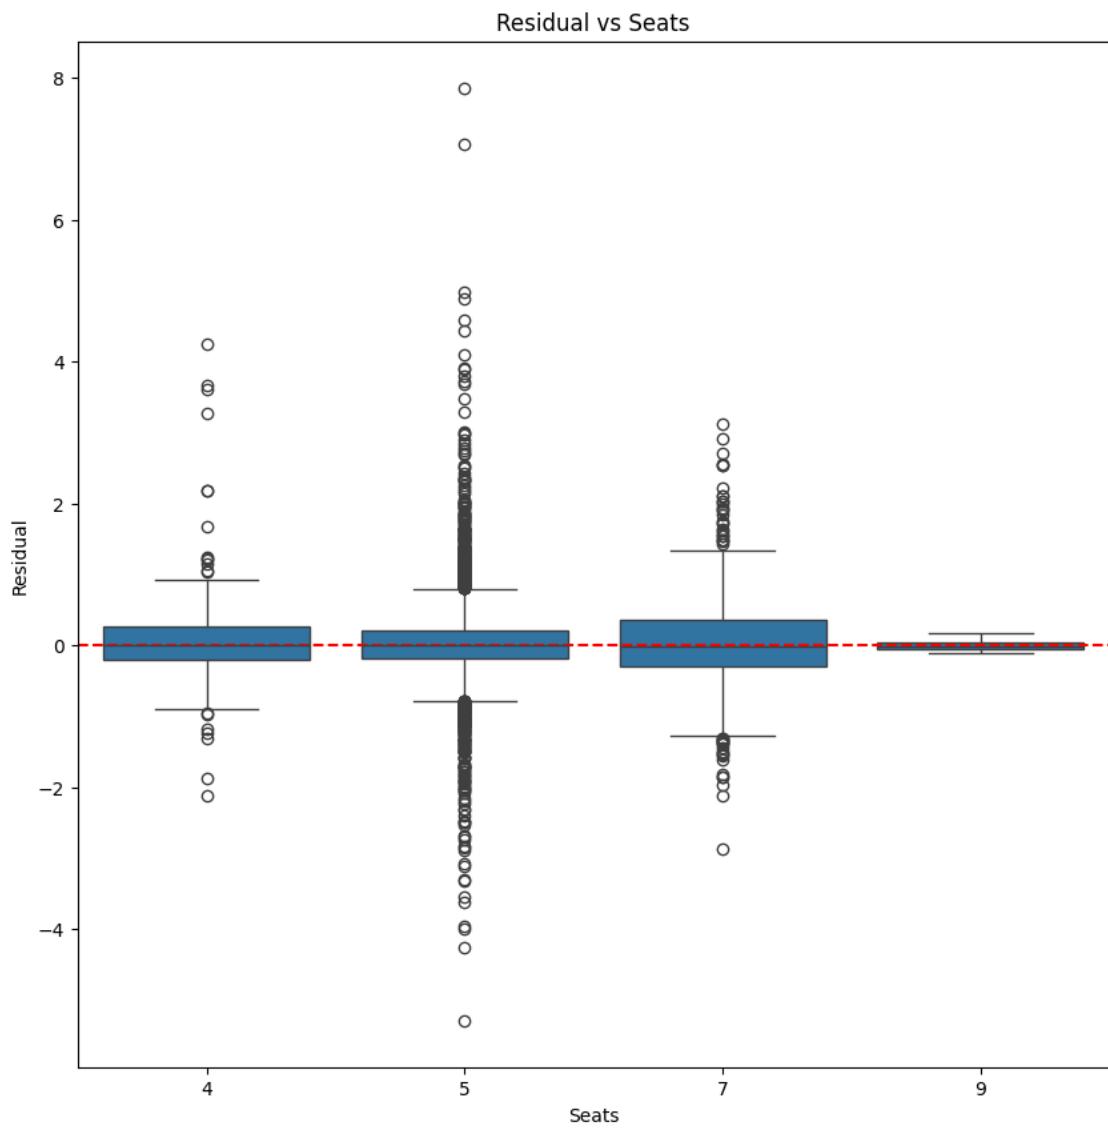


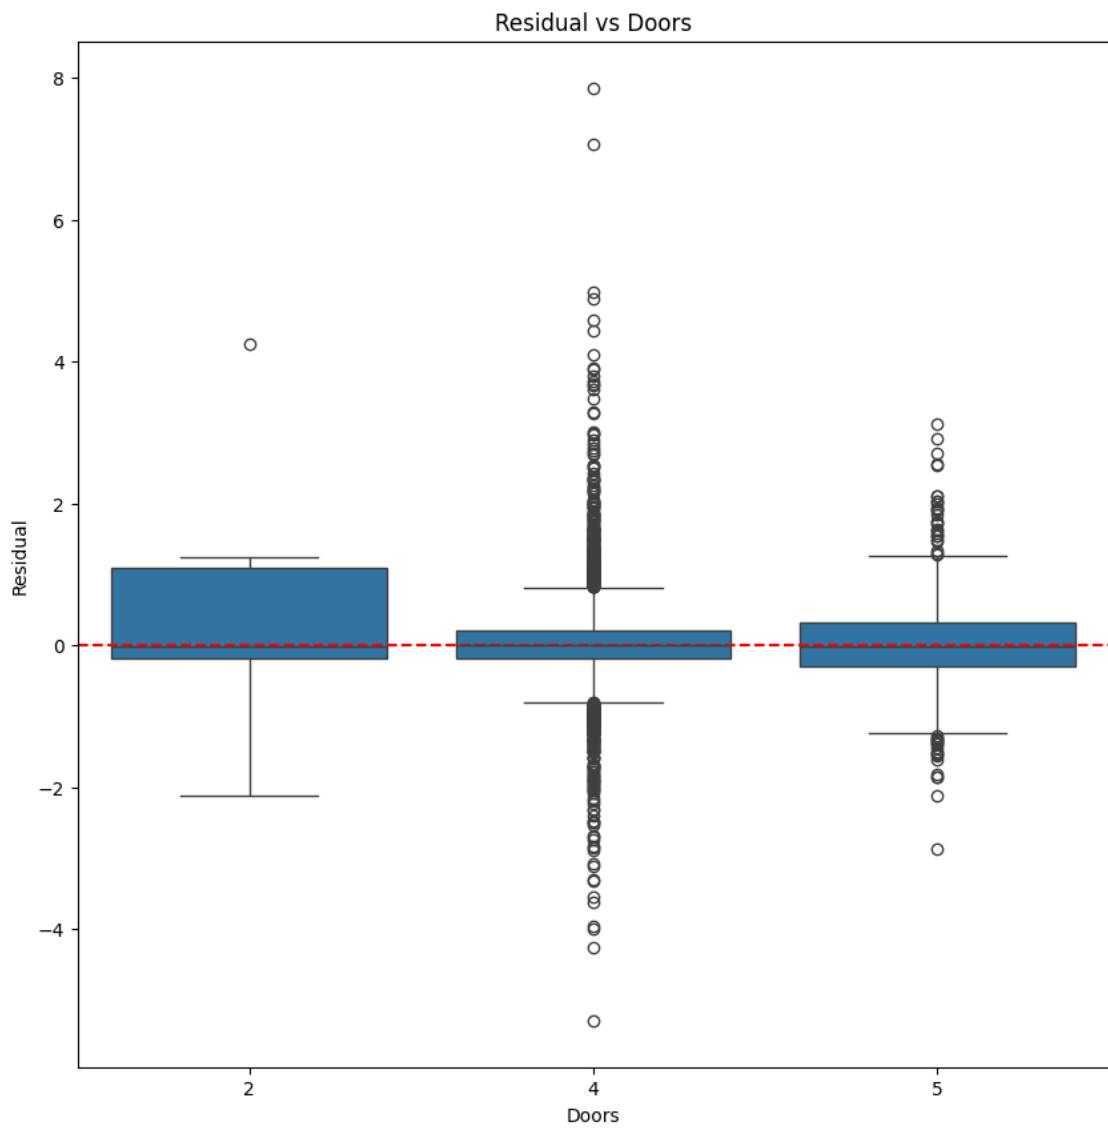
Residual vs Engine

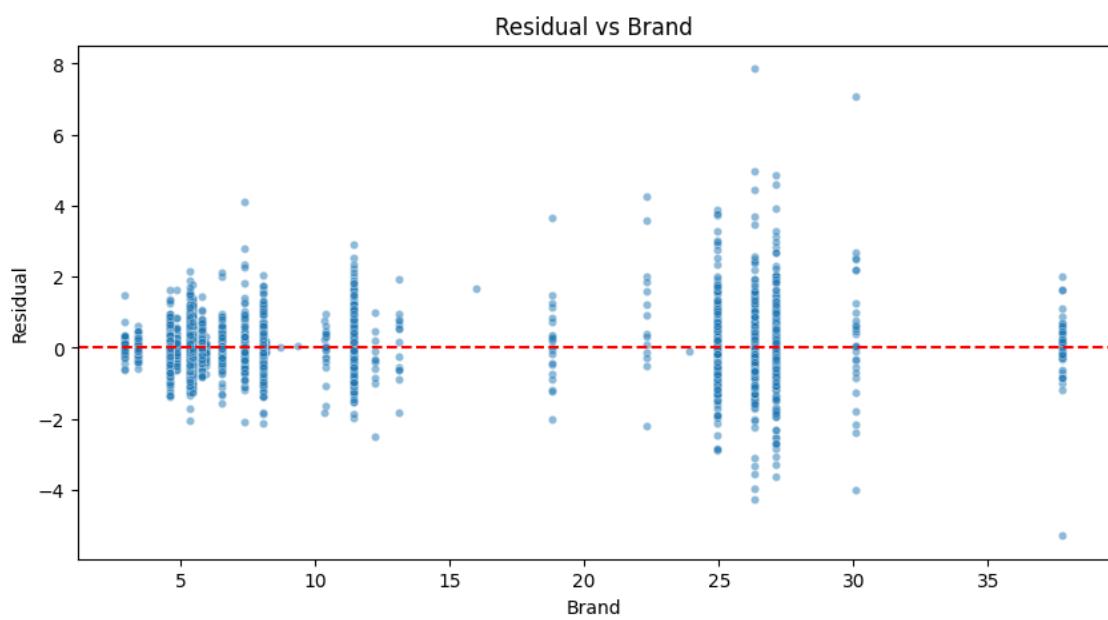
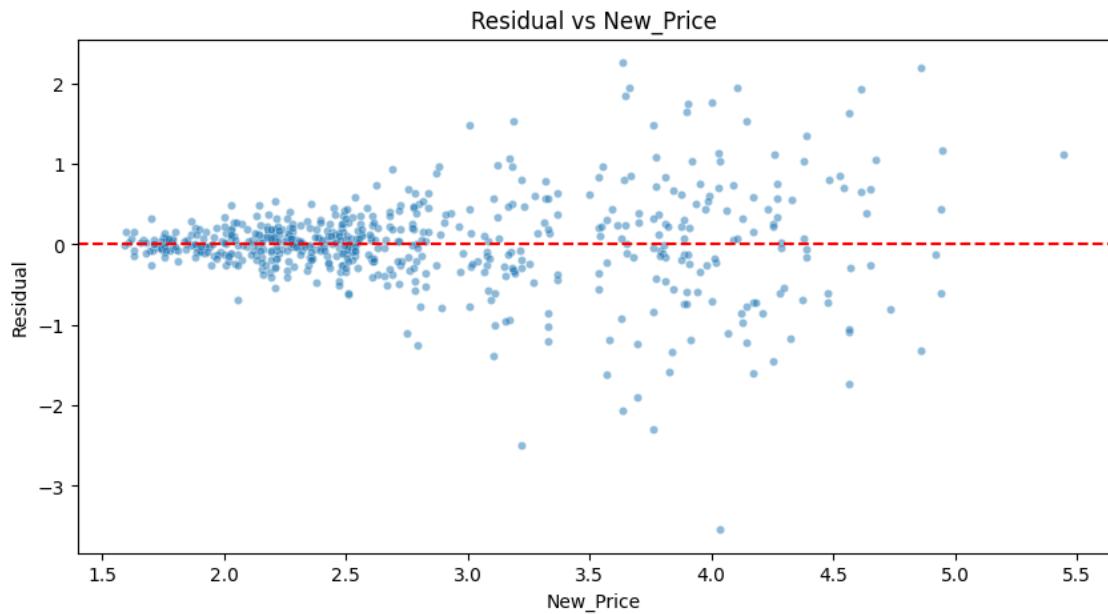


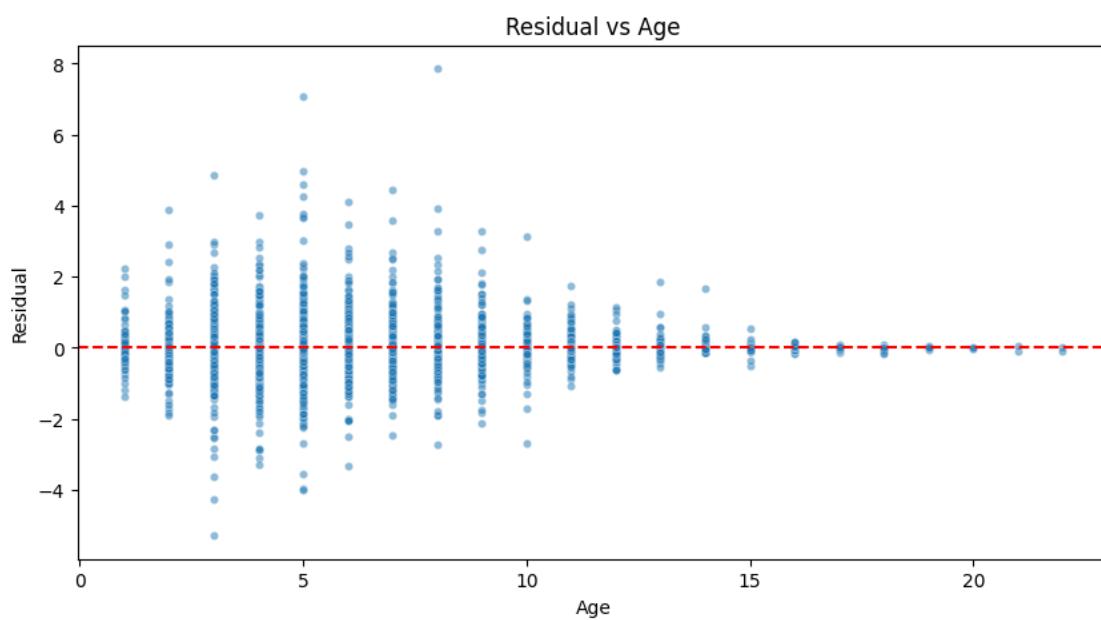
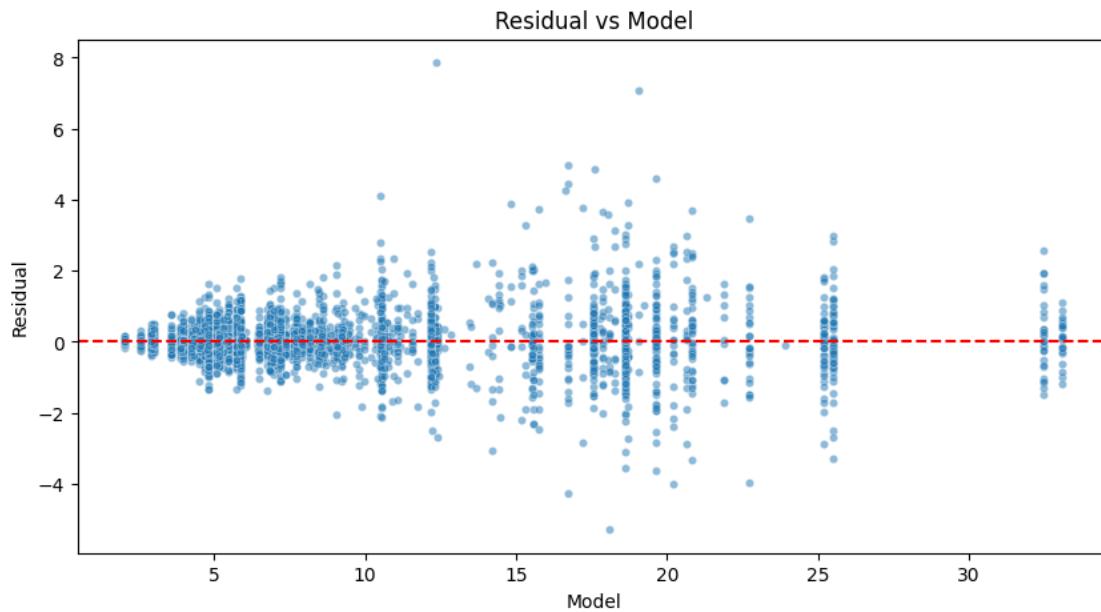


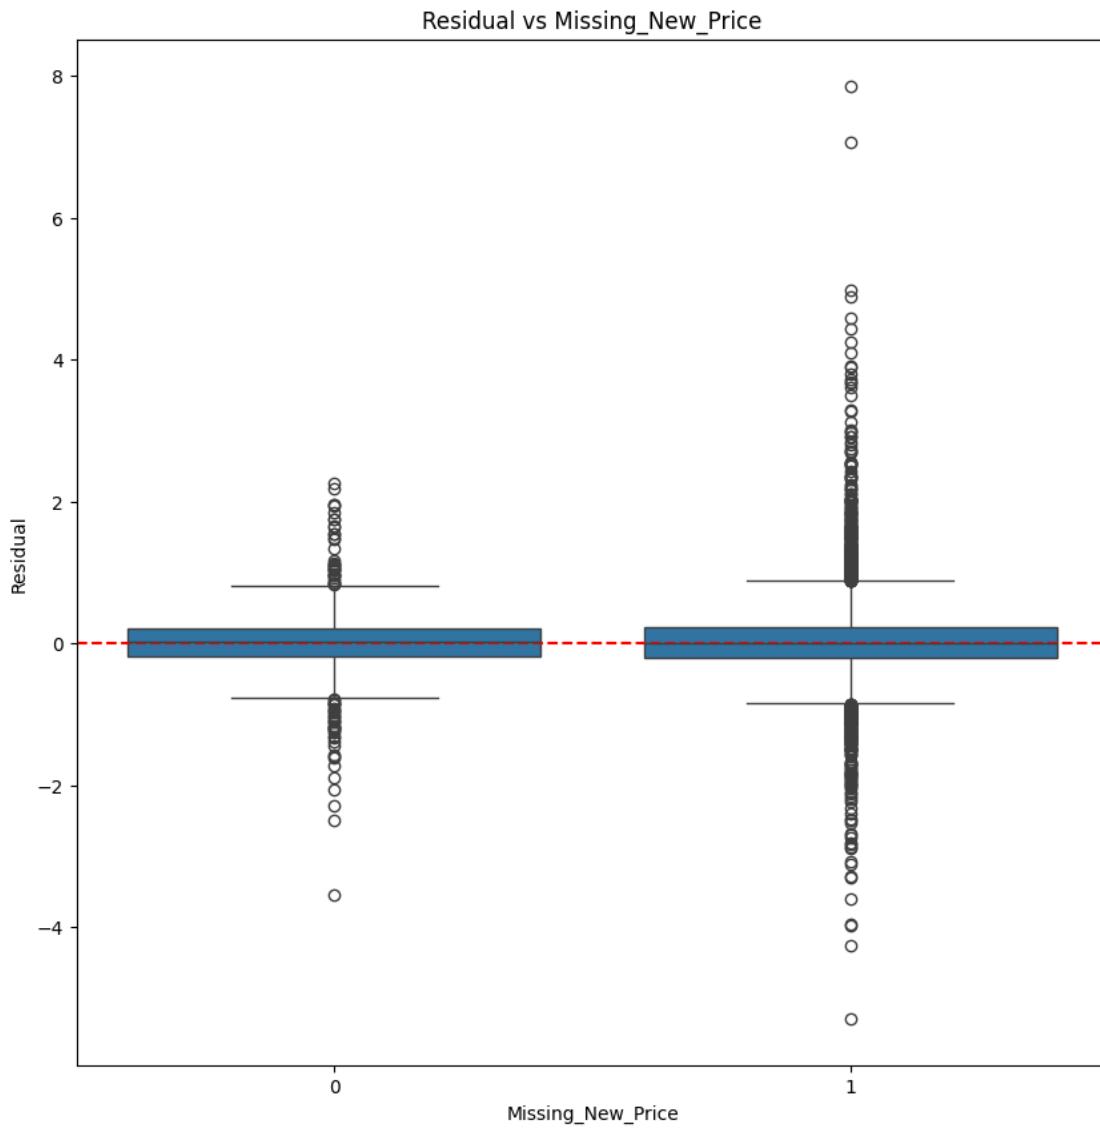






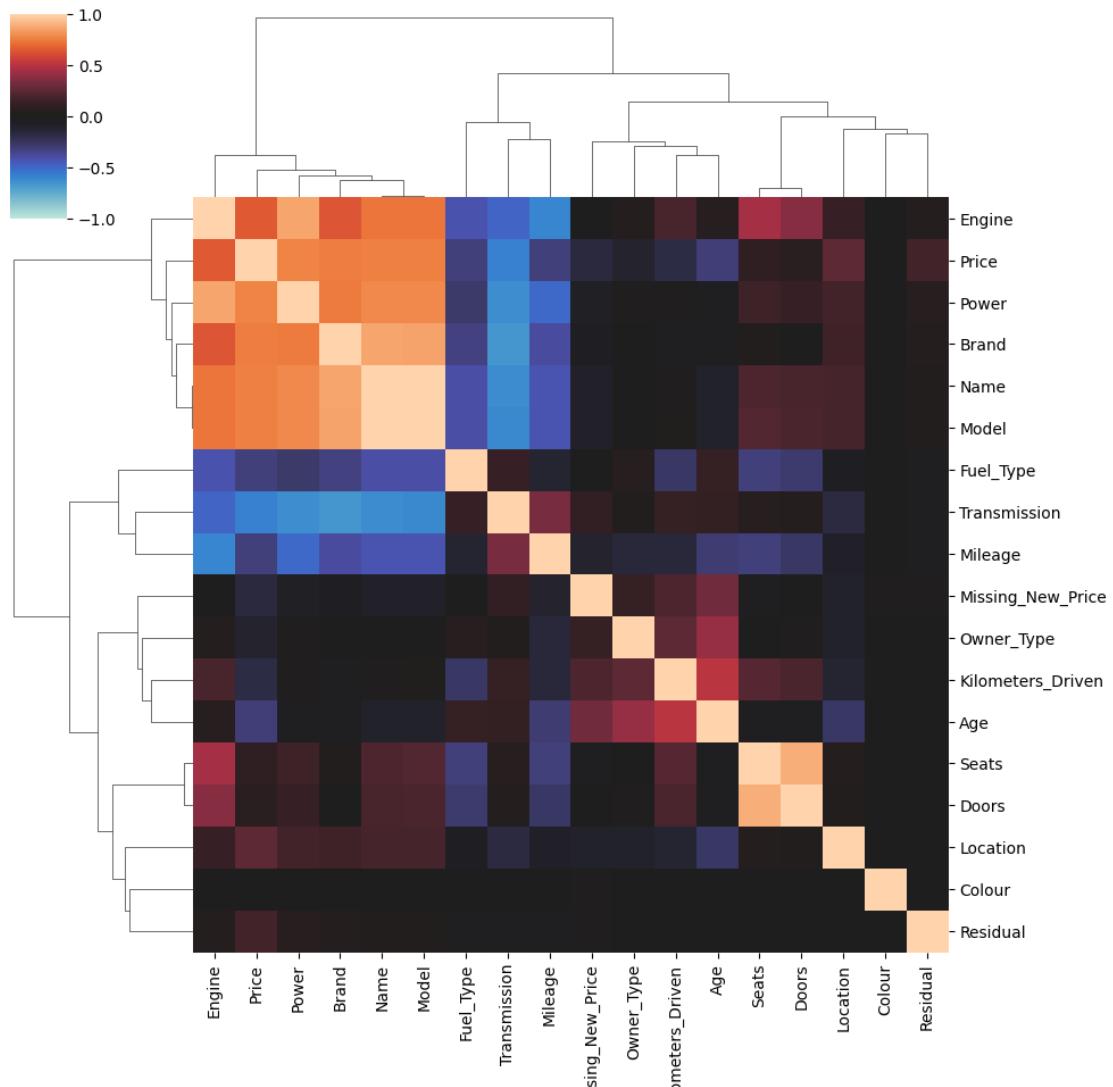






```
[88]: def plot_correlation(df, method="pearson", annot=True, **kwargs):
    sns.clustermap(
        df.corr(method, numeric_only=True),
        vmin=-1.0,
        vmax=1.0,
        cmap="icefire",
        method="complete",
        annot=annot,
        **kwargs,
    )
```

```
plot_correlation(df_residuals.drop(columns=["New_Price", "Prediction"]),
                 annot=False)
```



There is highly correlated block.

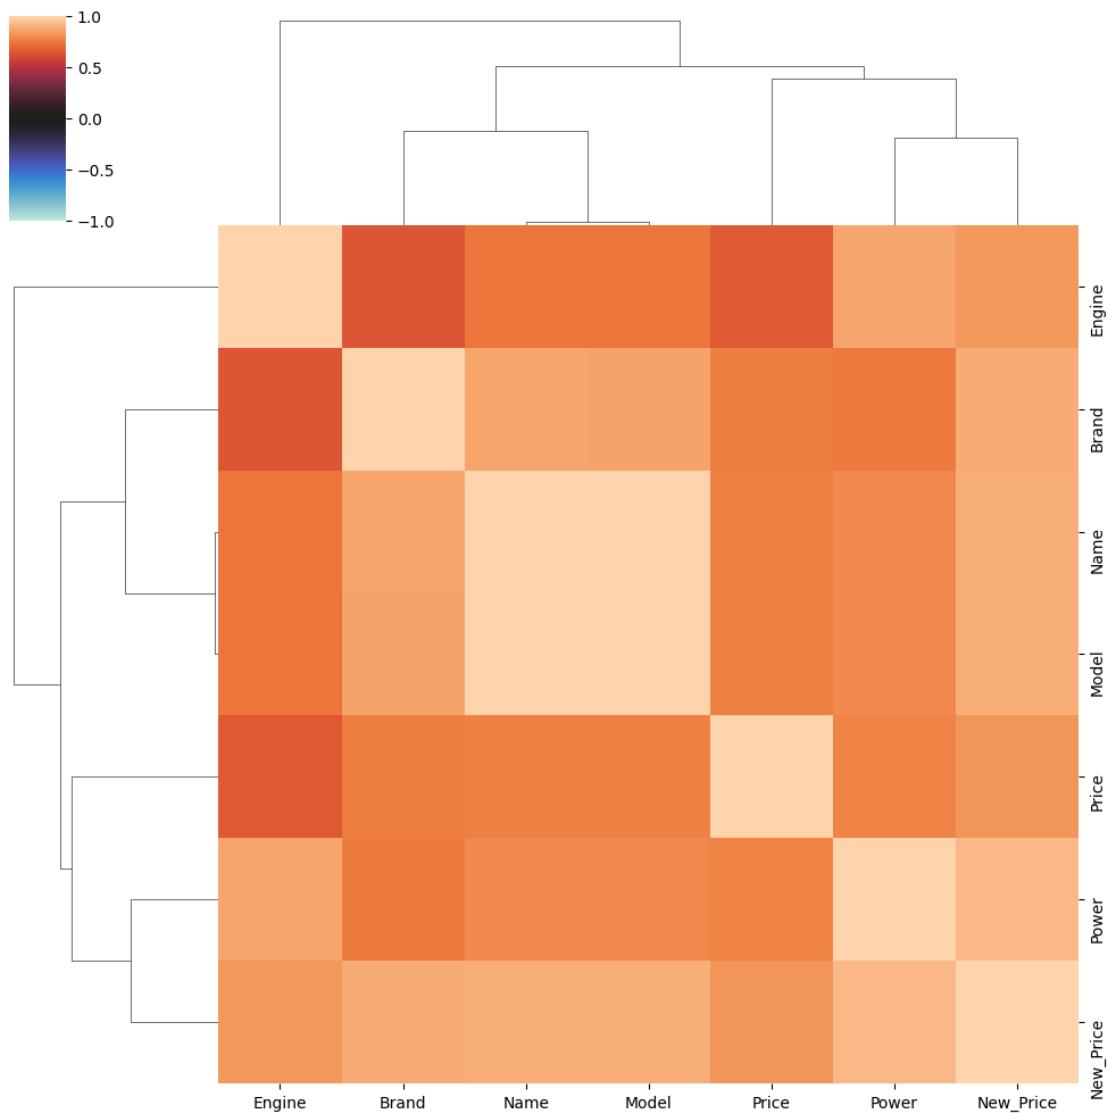
```
[89]: plot_correlation(
    df_residuals.drop(
        columns=[

            "Prediction",
            "Transmission",
            "Mileage",
            "Location",
            "Fuel_Type",
```

```

    "Seats",
    "Doors",
    "Missing_New_Price",
    "Owner_Type",
    "Kilometers_Driven",
    "Age",
    "Colour",
    "Residual",
]
),
annotate=False,
)

```



**Brand × Engine** The residual correlation matrix strongly suggests correlation between Brand and Engine. This is plausible since the Engine given a Brand can give a strong signal on its tier within the Brand.

```
[90]: class BrandEngineInteraction(TransformerMixin, BaseEstimator):
    def __init__(self,
                 brand_col: str = "Brand",
                 engine_col: str = "Engine",
                 brand_engine_col: str = "Brand-Engine",
                 ):
        self.brand_col = brand_col
        self.engine_col = engine_col
        self.brand_engine_col = brand_engine_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> Union["BrandEngineInteraction", None]:
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()
        X[self.brand_engine_col] = X[self.brand_col] * X[self.engine_col]
        return X
```

```
[91]: brand_engine_interaction_exp = Experiment(
    ExperimentConfig(
        name="brand_engine_interaction",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("brand_engine_interaction", BrandEngineInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

brand_engine_interaction_exp.run(
    X_train, y_train, None, best_feature_by_feature_exp_result
);
```

```
[Experiment: brand_engine_interaction]
Cross-validating (5-folds)...
CV score: 0.1382 ± 0.0161
    +0.0004  -0.0012 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete
```

It decreased the performance. This could be due to XGBoost already capturing such interaction implicitly, and the introduction of this interaction feature did not bring meaningfully more information.

**Brand × Power** Similarly, Brand and Power shows strong correlation and this is plausible with the same reason as why Brand and Engine interaction is meaningful.

```
[92]: class BrandPowerInteraction(TransformerMixin, BaseEstimator):
    def __init__(self,
                 brand_col: str = "Brand",
                 power_col: str = "Power",
                 brand_power_col: str = "Brand-Pwer",
                 ):
        self.brand_col = brand_col
        self.power_col = power_col
        self.brand_power_col = brand_power_col

    def fit(self, X: pd.DataFrame, y: Optional[pd.Series]) -> Union["BrandPowerInteraction", None]:
        return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        X = X.copy()
        X[self.brand_power_col] = X[self.brand_col] * X[self.power_col]
        return X
```

```
[93]: brand_power_interaction_exp = Experiment(
    ExperimentConfig(
        name="brand_power_interaction",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("brand_power_interaction", BrandPowerInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

brand_power_interaction_exp.run(
    X_train, y_train, None, best_feature_by_feature_exp_result
);
```

```
[Experiment: brand_power_interaction]
```

```
Cross-validating (5-folds)...
```

```

CV score: 0.1371 ± 0.0115
      -0.0008 -0.0058 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete

```

It improved the performance.

**Age × New\_Price** Age and New\_Price can represent depreciation. One of the most strongest model agnostic signal in predicting used-car price can be starting from New\_Price and monotonically decreasing as it Ages.

```
[94]: class AgeNewPriceInteraction(BaseEstimator, TransformerMixin):
    def __init__(self,
                 age_col="Age",
                 new_price_col="New_Price",
                 out_col="Age-NewPrice",
                 ):
        self.age_col = age_col
        self.new_price_col = new_price_col
        self.out_col = out_col

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        X = X.copy()
        X[self.out_col] = X[self.new_price_col] * X[self.age_col]
        return X
```

```
[95]: age_new_price_interaction_exp = Experiment(
    ExperimentConfig(
        name="age_new_price_interaction",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("age_new_price_interaction", AgeNewPriceInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

age_new_price_interaction_exp.run(
    X_train, y_train, None, best_feature_by_feature_exp_result
);
```

```
[Experiment: age_new_price_interaction]
Cross-validating (5-folds)...
CV score: 0.1368 ± 0.0159
    -0.0010 -0.0014 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete
```

It slightly increased the performance and stability.

```
[96]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("brand_power_interaction", BrandPowerInteraction()),
                ("age_new_price_interaction", AgeNewPriceInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

current_best_exp.run(X_train, y_train, None, ↴
    ↪best_feature_by_feature_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1380 ± 0.0105
    +0.0002 -0.0068 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete
```

I'll try more directly modeling depreciation.

The ratio of Price to New\_Price could tell the normalized amount of depreciation. Assuming constant rate of price drop per year.

```
[97]: X_train_intersection = X_train.dropna(subset=["New_Price"])
y_train_intersection = y_train.loc[X_train_intersection.index]
current_year = 2020

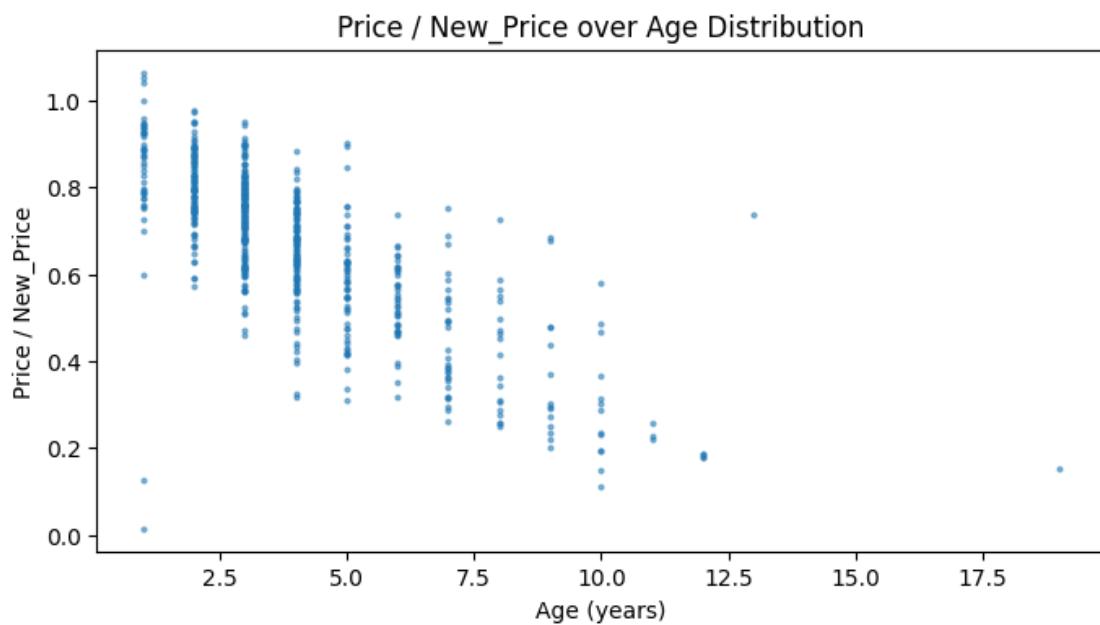
ages = current_year - X_train_intersection.Year
prices_over_new_price = y_train_intersection / X_train_intersection.New_Price
```

```

plt.figure(figsize=(8, 4))
plt.scatter(ages, prices_over_new_price, s=4, alpha=0.5)

plt.xlabel("Age (years)")
plt.ylabel("Price / New_Price")
plt.title("Price / New_Price over Age Distribution")
plt.show()

```



#### Model × Power

```

[98]: class ModelPowerInteraction(BaseEstimator, TransformerMixin):
    def __init__(
        self,
        model_col="Model",
        power_col="Power",
        model_power_col="Model-Power",
    ):
        self.model_col = model_col
        self.power_col = power_col
        self.model_power_col = model_power_col

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        X = X.copy()

```

```
X[self.model_power_col] = X[self.model_col] * X[self.power_col]
return X
```

```
[99]: model_power_interaction_exp = Experiment(
    ExperimentConfig(
        name="model_power_interaction",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("model_power_interaction", ModelPowerInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

model_power_interaction_exp.run(
    X_train, y_train, None, best_feature_by_feature_exp_result
);
```

```
[Experiment: model_power_interaction]
Cross-validating (5-folds)...
CV score: 0.1360 ± 0.0141
    -0.0018 -0.0032 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete
```

```
[100]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=Pipeline(
            [
                *best_single_feature_pipeline.steps[:-2],
                ("brand_power_interaction", BrandPowerInteraction()),
                ("model_power_interaction", ModelPowerInteraction()),
                *best_single_feature_pipeline.steps[-2:],
            ]
        ),
    )
)

current_best_exp.run(X_train, y_train, None,
    ↪best_feature_by_feature_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
```

```

CV score: 0.1372 ± 0.0116
      -0.0007 -0.0057 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete

```

## Ontology Approach

```
[101]: df_train.columns
```

```
[101]: Index(['Name', 'Location', 'Year', 'Kilometers_Driven', 'Fuel_Type',
       'Transmission', 'Owner_Type', 'Mileage', 'Engine', 'Power', 'Colour',
       'Seats', 'Doors', 'New_Price', 'Price'],
       dtype='object')
```

Purely data-driven approach is failing. Attempting to build a strongly domain-knowledge based ontology of features.

```
[102]: used_car_price_ontology = {
    "Car_Spec": {
        "Power": ["Engine", "Power", "Fuel_Type", "Transmission"],
        "Form": ["Seats", "Doors", "Colour"],
        "Positioning": ["Brand", "Model", "Name", "New_Price"],
    },
    "Condition": {"Wear": ["Age", "Kilometers_Driven"], "Ownership": [
        "Owner_Type"]},
    "Market": ["Location"],
}
```

The Car\_Spec will determine the car's BaseValue, the Condition will determine the car's Depreciation, and the Market will determine the MarketInfluence.

```
[103]: class BaseValueIndexTransformer(TransformerMixin, BaseEstimator):
    def __init__(
        self,
        brand_col="Brand",
        model_col="Model",
        fuel_col="Fuel_Type",
        trans_col="Transmission",
        engine_col="Engine",
        power_col="Power",
        seats_col="Seats",
        doors_col="Doors",
        name_col="Name",
        new_price_col="New_Price",
        n_segments=3,
    ):
        self.brand_col = brand_col
```

```

    self.model_col = model_col
    self.fuel_col = fuel_col
    self.trans_col = trans_col
    self.engine_col = engine_col
    self.power_col = power_col
    self.seats_col = seats_col
    self.doors_col = doors_col
    self.name_col = name_col
    self.new_price_col = new_price_col
    self.n_segments = n_segments

def _spec_features(self, df: pd.DataFrame) -> pd.DataFrame:
    out = pd.DataFrame(index=df.index)

    out["LogEngine"] = np.log1p(df[self.engine_col])
    out["LogPower"] = np.log1p(df[self.power_col])
    out["PowerDensity"] = df[self.power_col] / df[self.engine_col]

    out[self.seats_col] = df[self.seats_col]
    out[self.doors_col] = df[self.doors_col]

    return out

def fit(
    self, X: pd.DataFrame, y: Optional[pd.Series] = None
) -> "BaseValueIndexTransformer":
    df = X.copy()

    mask = df[self.new_price_col].notna()
    df_known = df.loc[mask].copy()

    num_df = self._spec_features(df_known)

    cat_df = df_known[
        [self.brand_col, self.model_col, self.fuel_col, self.trans_col]
    ]

    self.ohe_ = OneHotEncoder(handle_unknown="ignore", sparse_output=False)
    cat_arr = self.ohe_.fit_transform(cat_df)

    num_arr = num_df.fillna(0).values
    X_spec = np.hstack([cat_arr, num_arr])

    self.kmeans_ = KMeans(n_clusters=self.n_segments, random_state=42)
    segments = self.kmeans_.fit_predict(X_spec)
    df_known["Segment"] = segments

```

```

    self.num_features_ = num_df.columns.tolist()

    X_reg = np.hstack([X_spec, segments.reshape(-1, 1)])
    y_reg = df_known[self.new_price_col]

    self.reg_ = LinearRegression()
    self.reg_.fit(X_reg, y_reg)

    self.base_mean_ = y_reg.mean()

    return self

def transform(self, X: pd.DataFrame) -> pd.DataFrame:
    df = X.copy()

    num_df = self._spec_features(df)

    cat_df = df[[self.brand_col, self.model_col, self.fuel_col, self.
    ↪trans_col]]
    cat_arr = self.ohe_.transform(cat_df)

    num_arr = num_df.fillna(0).values

    X_spec = np.hstack([cat_arr, num_arr])

    segments = self.kmeans_.predict(X_spec)
    df["Segment"] = segments

    X_reg = np.hstack([X_spec, segments.reshape(-1, 1)])

    estimated_np = self.reg_.predict(X_reg)

    df["BaseValueIndex"] = estimated_np / self.base_mean_

    return df

```

```

[104]: class WearIndexTransformer(TransformerMixin, BaseEstimator):
    def __init__(
        self,
        age_col="Age",
        km_col="Kilometers_Driven",
        engine_col="Engine",
        owner_col="Owner_Type",
    ):
        self.age_col = age_col
        self.km_col = km_col
        self.engine_col = engine_col

```

```

    self.owner_col = owner_col

def _build_wear_features(self, df: pd.DataFrame) -> pd.DataFrame:
    out = pd.DataFrame(index=df.index)

    age = df[self.age_col].fillna(0)
    km = df[self.km_col].fillna(0)

    engine = df[self.engine_col].fillna(df[self.engine_col].median())

    out["Age_scaled"] = age
    out["KM_per_Year"] = km / (age + 1)
    out["WearFactor"] = km / (engine + 1)

    return out

def fit(
    self, X: pd.DataFrame, y: Optional[pd.Series] = None
) -> "WearIndexTransformer":
    df = X.copy()

    wear_df = self._build_wear_features(df)

    owner_vals = df[self.owner_col].astype("category").cat.codes
    wear_df["OwnerEncoded"] = owner_vals

    self.scaler_ = StandardScaler()
    wear_scaled = self.scaler_.fit_transform(wear_df)

    self.reg_ = LinearRegression()
    self.reg_.fit(wear_scaled, np.arange(len(wear_scaled)))

    self.w_ = self.reg_.coef_

    return self

def transform(self, X: pd.DataFrame) -> pd.DataFrame:
    df = X.copy()

    wear_df = self._build_wear_features(df)
    owner_vals = df[self.owner_col].astype("category").cat.codes
    wear_df["OwnerEncoded"] = owner_vals

    wear_scaled = self.scaler_.transform(wear_df)

    wear_score = wear_scaled @ self.w_

```

```

        wear_index = (wear_score - wear_score.min()) / (
            wear_score.max() - wear_score.min()
        )

        df["WearIndex"] = wear_index

    return df

```

```

[105]: class MarketIndexTransformer(TransformerMixin, BaseEstimator):
    def __init__(self, location_col="Location"):
        self.location_col = location_col

    def fit(
        self, X: pd.DataFrame, y: Optional[pd.Series] = None
    ) -> "MarketIndexTransformer":
        df = X.copy()
        df["Price"] = y

        loc_means = df.groupby(self.location_col)[("Price")].mean()

        df["LocMean"] = df[self.location_col].map(loc_means)

        self.ohe_ = OneHotEncoder(handle_unknown="ignore", sparse_output=False)
        loc_ohe = self.ohe_.fit_transform(df[[self.location_col]])

        self.reg_ = LinearRegression()
        self.reg_.fit(loc_ohe, df["LocMean"])

        self.global_mean_ = df["Price"].mean()

    return self

    def transform(self, X: pd.DataFrame) -> pd.DataFrame:
        df = X.copy()

        loc_ohe = self.ohe_.transform(df[[self.location_col]])

        loc_score = self.reg_.predict(loc_ohe)

        df["MarketIndex"] = loc_score / self.global_mean_

    return df

```

```

[106]: class OntologyFeatureBuilder(BaseEstimator, TransformerMixin):
    def __init__(
        self,
        base_transformer,

```

```

wear_transformer,
market_transformer,
add_interactions=True,
add_composite=True,
):
    self.base_transformer = base_transformer
    self.wear_transformer = wear_transformer
    self.market_transformer = market_transformer
    self.add_interactions = add_interactions
    self.add_composite = add_composite

def fit(self, X, y):
    self.base_transformer.fit(X)

    self.wear_transformer.fit(X)

    self.market_transformer.fit(X, y)

    return self

def transform(self, X):
    df = X.copy()

    df = self.base_transformer.transform(df)
    df = self.wear_transformer.transform(df)
    df = self.market_transformer.transform(df)

    base = df["BaseValueIndex"]
    wear = df["WearIndex"]
    market = df["MarketIndex"]

    if self.add_interactions:
        df["Base_x_Wear"] = base * wear
        df["Base_x_Market"] = base * market
        df["Wear_x_Market"] = wear * market

    if self.add_composite:
        df["CompositeValue"] = base * wear * market

    return df

```

[107]: ontology\_transformer = OntologyFeatureBuilder(
 base\_transformer=BaseValueIndexTransformer(),
 wear\_transformer=WearIndexTransformer(),
 market\_transformer=MarketIndexTransformer(),
 add\_interactions=True,
 add\_composite=True,

```
)
```

```
[108]: ontology_exp = Experiment(  
    ExperimentConfig(  
        name="ontology-based-engineering",  
        pipeline=Pipeline(  
            [  
                *best_single_feature_pipeline.steps[:-2],  
                ("ontology", ontology_transformer),  
                *best_single_feature_pipeline.steps[-2:],  
            ]  
        ),  
    ),  
)  
  
ontology_exp.run(X_train, y_train, None, best_feature_by_feature_exp_result);
```

```
[Experiment: ontology-based-engineering]  
Cross-validating (5-folds)...  
CV score: 0.1475 ± 0.0179  
    +0.0096  +0.0005 compared to combine-all-feature-by-feature-engineering  
(Negative is better)  
Training on full training set...  
Experiment complete
```

---

## Final Feature Engineering Pipeline

```
[109]: def final_pipeline_builder(model_config=dict()) -> Pipeline:  
    return Pipeline(  
        [  
            ("extract_brand_model", BrandModelExtractor()),  
            (  
                "target_encode",  
                TargetEncoder(cols=["Brand", "Model", "Name", "Location"]),  
            ),  
            ("transform", YearToAgeTransformer()),  
            ("Kilometers_Driven_clip_outliers", KilometersDrivenClipper()),  
            ("group_infrequent_fuel_type", FuelTypeGrouper()),  
            ("mileage_clip_outliers", MileageClipper()),  
            ("imput_power", PowerImputer()),  
            ("bin_seats", SeatsBinner()),  
            ("transform_new_price", NewPriceTransformer()),  
            ("brand_power_interaction", BrandPowerInteraction()),  
            ("model_power_interaction", ModelPowerInteraction()),  
            ("category_encode", CategoricalEncoder()),  
            (  
        ]
```

```
        "model",
        TransformedTargetRegressor(
            regressor=XGBRegressor(**model_config),
            func=np.log1p,
            inverse_func=np.expm1,
        ),
    ),
],
)
```

```
[110]: current_best_exp = Experiment(
    ExperimentConfig(
        name="current-best",
        pipeline=final_pipeline_builder(),
    )
)

current_best_exp.run(X_train, y_train, None, ↴
    best_feature_by_feature_exp_result);
```

```
[Experiment: current-best]
Cross-validating (5-folds)...
CV score: 0.1372 ± 0.0116
    -0.0007  -0.0057 compared to combine-all-feature-by-feature-engineering
(Negative is better)
Training on full training set...
Experiment complete
```