

hyperparameter-tuning

November 17, 2025

0.0.1 Hyperparameter Tuning

```
[1]: %load_ext autoreload
%autoreload 2

[2]: from mldl_hw3.preprocessing import DataLoader
from mldl_hw3.feature_engineering import build_feature_engineering_pipeline
from mldl_hw3.experiment import Experiment, ExperimentConfig

from itertools import product

from xgboost import XGBRegressor
from tqdm import tqdm
from IPython.display import clear_output
import pandas as pd

[3]: df_train, df_test = DataLoader("../dataset").load()

X_train = df_train.copy()
y_train = X_train.pop("Price")
X_test = df_test.drop(columns=["Price"])
```

Grid Search

```
[4]: grid_config = {
    "max_depth": [4, 6],
    "min_child_weight": [1, 5],
    "gamma": [0.0, 0.1],
    "reg_lambda": [1, 2, 5],
    "reg_alpha": [0, 0.1],
    "subsample": [0.7, 0.9],
    "colsample_bytree": [0.7, 0.9],
    "learning_rate": [0.03, 0.1],
    "n_estimators": [800, 1500],
    "objective": ["reg:squarederror"],
    "tree_method": ["hist"],
    "random_state": [42],
}
```

```

grid = [
    dict(zip(grid_config.keys(), combination))
    for combination in product(*grid_config.values())
]

```

```

[5]: grid_search_results = []

for i, params in tqdm(enumerate(grid), total=len(grid)):
    exp = Experiment(
        ExperimentConfig(
            name=f"xgb-params-grid-search-{i}",
            pipeline=build_feature_engineering_pipeline(XGBRegressor(**params)),
            extra={"xgb-params": params},
        )
    )

    exp_result = exp.run(X_train, y_train, skip_full_training=True)

    clear_output(wait=True)
    grid_search_results.append((params, exp_result))

```

100% | 768/768 [09:35<00:00, 1.33it/s]

```

[6]: param_keys = grid_search_results[0][0].keys()

df_grid_search_results = pd.DataFrame(
{
    **{key: [d[key] for d, _ in grid_search_results] for key in param_keys},
    "MAPE": [exp_result.cv_score for _, exp_result in grid_search_results],
    "MAPE_std": [exp_result.cv_std for _, exp_result in
    grid_search_results],
}
).sort_values(by=["MAPE"])
df_grid_search_results

```

	max_depth	min_child_weight	gamma	reg_lambda	reg_alpha	subsample	\
46	4		1	0.0	2	0.0	0.9
1	4		1	0.0	1	0.0	0.7
47	4		1	0.0	2	0.0	0.9
34	4		1	0.0	2	0.0	0.7
35	4		1	0.0	2	0.0	0.7
..	
312	4		5	0.1	1	0.1	0.9
120	4		1	0.1	1	0.1	0.9
570	6		1	0.1	5	0.1	0.9
376	4		5	0.1	5	0.1	0.9
360	4		5	0.1	5	0.0	0.9

```

  colsample_bytree  learning_rate  n_estimators      objective \
46            0.9          0.10        800  reg:squarederror
1             0.7          0.03       1500  reg:squarederror
47            0.9          0.10       1500  reg:squarederror
34            0.7          0.10        800  reg:squarederror
35            0.7          0.10       1500  reg:squarederror
..
312           0.7          0.03        800  reg:squarederror
120           0.7          0.03        800  reg:squarederror
570           0.7          0.10        800  reg:squarederror
376           0.7          0.03        800  reg:squarederror
360           0.7          0.03        800  reg:squarederror

  tree_method  random_state      MAPE   MAPE_std
46         hist           42  0.126625  0.015765
1          hist           42  0.127095  0.017024
47         hist           42  0.127118  0.016073
34         hist           42  0.127166  0.014453
35         hist           42  0.127531  0.014594
..
312        hist           42  0.149529  0.015915
120        hist           42  0.149531  0.016226
570        hist           42  0.149603  0.014402
376        hist           42  0.149861  0.015222
360        hist           42  0.150076  0.015829

```

[768 rows x 14 columns]

```
[7]: best_config = (
    df_grid_search_results.drop(columns=["MAPE", "MAPE_std"]).iloc[0].to_dict()
)
best_config
```

```
[7]: {'max_depth': 4,
 'min_child_weight': 1,
 'gamma': 0.0,
 'reg_lambda': 2,
 'reg_alpha': 0.0,
 'subsample': 0.9,
 'colsample_bytree': 0.9,
 'learning_rate': 0.1,
 'n_estimators': 800,
 'objective': 'reg:squarederror',
 'tree_method': 'hist',
 'random_state': 42}
```

```
[8]: xgb_pipeline = build_feature_engineering_pipeline(XGBRegressor(**best_config))
xgb_pipeline.fit(X_train, y_train)
test_predictions = xgb_pipeline.predict(X_test)
pd.DataFrame({"ID": X_test.index, "Price": test_predictions}).to_csv(
    "./artifacts/experiment-results/grid-search-tuning.csv", index=False
)
```